Detecting and Mitigating DDoS Attacks with AI: A Survey

ALEXANDRU APOSTU*, University of Bucharest, Romania

SILVIU GHEORGHE*, University of Bucharest, Romania

ANDREI HÎJI*, University of Bucharest, Romania

NICOLAE CLEJU, University of Bucharest, Romania

ANDREI PĂTRAȘCU, University of Bucharest, Romania

CRISTIAN RUSU, University of Bucharest, Romania

RADU TUDOR IONESCU, University of Bucharest, Romania

PAUL IROFTI†, University of Bucharest, Romania

Distributed Denial of Service attacks represent an active cybersecurity research problem. Recent research shifted from static rule-based defenses towards AI-based detection and mitigation. This comprehensive survey covers several key topics. Preeminently, state-of-the-art AI detection methods are discussed. An in-depth taxonomy based on manual expert hierarchies and an AI-generated dendrogram are provided, thus settling DDoS categorization ambiguities. An important discussion on available datasets follows, covering data format options and their role in training AI detection methods together with adversarial training and examples augmentation. Beyond detection, AI based mitigation techniques are surveyed as well. Finally, multiple open research directions are proposed.

CCS Concepts: • **Networks → Network monitoring**; • **Security and privacy → Intrusion/anomaly detection and malware mitigation**; **Intrusion detection systems**; **Denial-of-service attacks**; • **Computing methodologies → Machine learning**; **Artificial intelligence**.

Additional Key Words and Phrases: DDoS, DoS, DDoS Detection, DDoS Mitigation, Artificial Intelligence, Machine Learning

---

*Authors contributed equally to this research.
†Corresponding author.

---

1 Introduction

DoS (Denial of Service) and DDoS (Distributed Denial of Service) represent a major cybersecurity threat. DDoS attacks increased by 55% between January 2020 and March 2021 [95], a phenomenon that was accentuated starting from the second half of 2021 and continued until the first half of 2022, when the total number of attacks increased by 75.60% [67]. According to Markets&Markets[1], the global DDoS security and protection market size was valued at USD 3.9 billion in 2022 and is expected to reach USD 7.3 billion by 2027, with a CAGR of 13.2% from 2022 to 2027. Cloud-based DDoS security and protection services can effectively handle volumetric DDoS attacks, as well as layer 3 and 7 attacks. Thus, to optimize operations and costs, companies are rapidly adopting cloud-based anti-DDoS solutions. At the same time, due to the increasing adoption of 5G technology, a survey by the American mobile operator A10 showed that 63% of respondents believe that advanced DDoS security and protection solutions are needed to protect 5G networks. Therefore, we consider DDoS detection and mitigation an important global topic where modern AI techniques are necessary to fight off attacks, while, at the same time, helping the victim in maintaining a functioning infrastructure.

In this work, we use machine learning and artificial intelligence interchangeably to describe learning-based methods. Where appropriate, we will use the terms shallow and deep learning to separate classical machine learning methods from deep neural networks. We follow the classification of DDoS attacks according to the Cybersecurity and Infrastructure Security Agency guide provided by the Federal Bureau of Investigation [26]. There, DDoS categories are separated into volumetric [70], protocol [70, 91], and application [70, 91, 104] level attacks [97]. Volumetric attacks are most common, often involving bot networks that bombard the victim with a large number of connections that do not have to maintain state, like UDP packets, for example. Protocol attacks are more sophisticated, as they exploit defects found in specific protocol implementations to employ both network floods, but also computational burden on the target. Reflection and amplification [70, 81] are commonly-used DDoS techniques meant to widen volumetric and protocol attacks. Application attacks target high-level services, such as web applications and database management systems, and mostly inflict large computational loads on the victims.

DDoS defenses are of special interest for IoT networks which are more susceptible and represent a large part of the bot infrastructure mounted during volumetric and protocol attacks [45, 75]. SDN is another particular context where anti-DDoS methodologies were studied [66, 98]. As these represent dynamic network architectures often used in cloud environments, their configuration usually includes recent DDoS defenses which take advantage of the known software-defined topology.

AI-based detection is often treated in the literature as a binary classification task [83, 88] performed on available datasets [9, 92, 93] containing raw labeled traffic recordings of normal and attack traffic. The attacks are usually mixed [46, 64] as the datasets are designed for IDS [92] or IoT networks [8, 49], and not specifically for the DDoS case. That is why, among the datasets that do contain DDoS attacks, most focus on volumetric floods or web application attacks. Recent trends show an interest in the use of adversarial training and adversarial attacks and examples [5, 38] as a means to provide more robust detection models.

Once detected, the next critical step for anti-DDoS methods is attack mitigation [5]. Mitigation represents the set of measures taken to block the attack and permit normal traffic to pass through. These often resume to manual firewall rules, custom made for the current attack. Unfortunately, AI literature on the topic is limited to DTs [23, 121] and a few specific cases, such as DNS Floods [39, 112]. Only recently, we started to see an interest in specializing deep learning methods in general, and LLM in particular, to automatically generate firewall rules [57, 111, 119].

---

[1]https://www.marketsandmarkets.com/Market-Reports/ddos-protection-mitigation-market-111952874.html

Contributions. In this paper, we present a clear and structured survey covering all attack categories and attack types, with a special focus on their detection and mitigation through machine learning methods. To our knowledge, this undertaking has not been addressed in the literature, so far. As such, we would like to underline the following important contributions that our study provides:

(1) *Taxonomy.* (a) We provide a thorough hierarchical grouping of the surveyed research papers from the perspective of their main contribution; (b) we identify four different attack categories in an attempt to settle existing overlaps and ambiguity found in the literature, especially between protocol and application attacks; (c) we introduce an automatic taxonomy based on an agglomerative clustering algorithm, which we discuss and compare;

(2) *Data format.* As far as we are aware, our work is the first to discuss and analyze upfront data formatting and preprocessing options for the machine learning algorithms; starting from the raw network traffic, we discuss data representation as topological graph data for traffic structure information, timeseries data for aggregate traffic evolution, and tabular data for individual packets and coalesced flow information;

(3) *Detection.* (a) We provide an in-depth discussion for each of the four attack categories including all attack types within; (b) we include an important discussion and analysis regarding the time delay between the moment that an attack is mounted and the moment when it is detected and then mitigated;

(4) *Mitigation.* As far as we are aware, we are the first to survey papers dealing with mitigation, after the DDoS attack is detected, where mitigation is achieved trough AI methods whose task is to emit efficient firewall blocking rules;

(5) *AI-generated DDoS traffic.* This is the first survey to include the topic of (a) adversarial training and (b) adversarial examples and attacks – these generative AI techniques can help to improve detection and mitigation;

(6) *Research directions.* We identify and provide clear and direct ways of improving the field of DDoS detection and mitigation using AI; although the literature consists of a vast amount of work on this topic, we notice a lack of specialized anti-DDoS machine learning methods, while practical tasks, such as minimizing detection delays together with providing efficient mitigation techniques, are just beginning to be explored.

Outline. The survey is organized in multiple sections with the list of terms and abbreviations collected in Table 1. Section 2 discusses existing surveys by comparing our work with them and motivating the need for the current manuscript. In Section 3, we carefully define and describe the attack categories and attack types within each category in order to group the selected papers from this survey through the use of both manual and automatic taxonomies. Section 4 discusses available datasets and, more importantly for machine learning algorithms, the data formats and preprocessing techniques that are used by recent learning algorithms. Next, we delve into the details regarding DDoS detection, in Section 5, and AI-based mitigation, in Section 7. The detection section includes all the attack categories included in the taxonomy (volumetric, protocol and application), but also a dedicated subsection for reflection and amplification techniques. In Section 6, we discuss how to further improve existing algorithms, models and datasets through adversarial training and adversarial examples. Finally, in Section 8, we discuss future research opportunities for the directions described above along with conclusion after our intensive survey activity.

2 Related Surveys

While existing literature contains plenty survey papers on the DDoS topic, most are centered on the attacks and few on the detection and mitigation tasks. Out of the latter, even fewer include learning based approaches, as can be seen in

Acronym	Description	Acronym	Description
AS	Autonomous System	EAD	Elastic-net Attacks to DNNs [20]
BGP	Border Gateway Protocol	ED	Euclidean distance
DDoS	Distributed Denial of Service	EWMA	Exponentially Weighted Moving Average [99]
DoS	Denial of Service	FCM	Fuzzy C-means [72]
DrDoS	Distributed reflection DoS	FGSM	Fast Gradient Sign Method [85]
HDFD	HTTP DDoS Flooding Defender	FL	Fuzzy Logic [65]
IDF	Inverse Document Frequency	GA	Genetic Attack [41]
IDS	IDS	GAN	Generative Adversarial Networks [123]
IoT	Internet of Things	GMM	Gaussian Mixtures Model [15]
MAN	Metropolitan Area Networks	GNB	Gaussian Naive-Bayes [82]
PoD	Ping of Death	GRU	Gated Recurrent Units [21]
RA-DDoS	Reflection and amplification DDoS	IF	Isolation Forest [117]
R.U.D.Y.	R U Dead Yet attack	JSMA	Jacobian-based Saliency Map Attack [115]
SDN	Software Defined Network	k-NN	k-Nearest Neighbors [25]
SIP	Session Initiation Protocol	LGBM	Light Gradient Boosting Machine [33]
TF	Term Frequency	LLM	Large Language Model [126]
TF-IDF	TF-Inverse Document Frequency	LoRA	Low-Rank Adaptation [116]
ToS	Type of Service	LSTM	Long Short-Term Memory [106]
TTL	Time To Live	MLP	Multi-Layer Perceptron
XDP	eXpress DataPath	MNB	Multinomial Naive-Bayes [82]
AIC	Akaike Information Criterion [18]	NB	Naive Bayes [82]
ANOVA	Analysis of variance [78]	OC-SVM	One-Class SVM [90]
BiGAN	Bidirectional GAN [124]	PEFT	Parameter Efficient Fine-Tuning [36]
CNN	Convolutional Neural Networks [42]	PWPSA	Probability Weighted Packet Saliency Attack [41]
CoD	Constraint-of-Deviation [55]	RF	Random Forest
CoT	Chain-of-Thought [120]	SVM	Support Vector Machine
CTGAN	Conditional GAN [16]	TVAE	Tabular Variational Autoencoder [118]
DAGMM	Deep Autoencoding GMM [127]	VAE	Variational Autoencoder [34]
DT	Decision Trees [24]	XGB	Extreme Gradient Boosting

Table 1. List of terms and acronyms used in the survey. Grouped at the top are network and attack terms in alphabetical order, followed at the bottom and to the right by artificial intelligence nomenclature. For the interested reader, we added references to recent surveys and papers that discuss the topic further.

Table 2: only the three papers at the top employ a general study of the field, while the ones at the bottom focus on specific tasks and testbeds.

Anti-DDoS surveys studying the general problem from a machine learning perspective focus exclusively on the detection task, without discussing learning-based mitigations. Also, these generic surveys discuss only a share of existing DDoS categories, mostly involving a subset of volumetric, protocol or application attacks, but never all of them together. Malliga et al. [91] cover mainly protocol-based attacks consisting of a table of 66 research papers on the use of deep machine learning methods for DDoS detection and a table of 12 of the most popular datasets used in simulations. The article gives a good tabular description of the field, but does not detail the methods analyzed and does not discuss the methods of generating DDoS attacks. Volumetric attack detection and mitigation strategies are completely missing. Interested readers can continue with application-level attacks from the work of Tripathi et al. [104], who provide a thorough survey regarding exclusively these types of attacks and detection mechanisms, while also including existing turnkey solutions from the industry. Unfortunately, the authors tackle few existing learning-based detection techniques,

Survey	Volumetric		Protocol				Application			Generative		Mitigation	Limitations
	Flood	Reflection	TCP	Ping	BGP	Smurf	HTTP	DNS	Slow	Training	Adversarial		
Ours	✔✔	✔✔	✔✔	✔✔	✔✔	✔✔	✔✔	✔✔	✔✔	✔✔	✔✔	✔✔	
Malliga et al. [91]			✔✔	✔✔			✔✔			✔✔			Missing volumetric
Najafimehr et al. [70]	✔	✔	✔	✔		✔	✔	✔	✔				Missing generative
Tripathi et al. [104]							✔✔	✔✔	✔✔				Application only
Singh et al. [97]			✔			✔	✔		✔				Attack only
Nuiaa et al. [81]		✔						✔✔	✔✔				DRDoS and Attack only
Kadri et al. [45]	✔✔		✔✔							✔			Missing application, IoT only
Pakmehr et al. [75]	✔		✔	✔		✔	✔	✔	✔	✔			IoT only
Musa et al. [66]	✔✔		✔✔	✔✔			✔✔	✔	✔✔		✔✔		SDN only
Su et al. [98]	✔✔		✔✔		✔								Missing application, SDN only
Alatwi et al. [5]			✔✔							✔	✔✔	✔	Protocol only, Generative only
He et al. [38]							✔			✔✔	✔✔		Application only, Generative only

Table 2. Comparison of covered AI-based DDoS detection and mitigation topics in related surveys; ✔✔ means AI detection or mitigation is directly addressed in the survey, ✔ means the topic is mentioned but treated in bulk with others or indirectly covered.

most of which are shallow. Further on, Najafimehr et al. [70] present a detailed taxonomy of DDoS attack detection methods which includes the missing volumetric attack types. The survey is centered around pairings between general detection methods and existing databases with a bias towards shallow learning approaches. Detecting specific attacks is not treated as a separate topic nor the particularities of each task; the algorithms are expected to detect with high accuracy any type of DDoS attack present in the dataset.

It is worth mentioning that some of the DDoS attack surveys also include a small section dedicated to learning-based detection of these attacks. We mention here the work of Singh et al. [97], which enumerates the list of machine and deep learning papers dealing with some protocol and application based attacks, and that of Nuiaa et al. [81], which analyzes solely DDoS attacks against the DNS protocol focusing on distributed methods with reflection (and amplification) that includes a small discussion about using AI for the attack detection task.

Moving on towards specialized surveys, we note that reflection, amplification, mitigation and attack generation techniques are not discussed in these works. The IoT domain is presented by Kadri et al. [45], who discuss floods and TCP-based attacks analyzing detection and validation methods used in the literature. The paper talks little about the performance of these methods in simulations and practical applications. The work of Pakmehr et al. [75] picks-up on this and presents several tables with advantages and disadvantages for shallow and deep learning methods; these tables are not very well connected to the survey text or the particularities of the attacks exploited by the detection methods. The paper covers a wider area of attacks, including a brief mention of modern GAN-based learning strategies. Nonetheless, techniques for generating AI-based DDoS attacks are not discussed.

SDNs are a hot topic in the DDoS field, and indeed, we also find two attack detection surveys specializing on it. The work of Musa et al. [66] tackles the detection task with an in-depth and wide coverage across the DDoS attack classes describing recent learning methods (both classical and deep) including Adversarial AI approaches. The article laments the simplistic scenarios discussed in the research articles and the need to test more complex scenarios that integrate with already existing security systems. While the authors provide a thorough review of the SDN literature on the topic, the survey does not provide insights and connections between the selected papers employing instead a list of abstracts

approach. The mitigation task is not covered here nor in other SDN surveys. Even though Su et al. [98] conduct a study of most mitigation methods for SDN DDoS attacks, learning methods are used solely in the detection process. These range from statistical to shallow and deep learning methods together with hybrid variants.

While significant progress has been made in the field, there are still some major obstacles, especially regarding the quality of the datasets and the adaptability of the defense systems. Even though not DDoS-centric, Alatwi et al. [5] focus on three main issues: classifying studies from the existing literature that deal with generating adversarial examples, evaluating deep learning-based detection systems, and proposing defense mechanisms against such attacks. The authors state that approaches such as adversarial training, although effective in areas like computer vision, fail to perform as well in network IDS because of the complex structure of traffic characteristics. While making some interesting observations, the survey of Alatwi et al. [5] covers a narrow domain, adversarial learning, dealing with a generic problem, network intrusion detection.

Supported by the views in [5], He et al. [38] talk about the vulnerabilities and limitations of machine learning-based methods in network IDS in the context of adversarial attacks. As noted by previous research, these systems are vulnerable to attacks that manipulate network traffic to avoid detection. Still, He et al. [38] argue that many of these attacks lack adaptability to the data structure of the network. They conclude that the literature still needs more representative and diverse datasets, along with more robust defense mechanisms such as adversarial training and feature reduction strategies. The survey of He et al. [38] is focused on a specific area related to adversarial attack generation.

3 Taxonomies of DDoS Attacks and AI-Based Detection/Mitigation Solutions

We start by introducing a set of clear definitions for the DDoS attack categories that are meant to easily and unambiguously separate existing attack types. Next, we provide a manual hierarchical grouping (taxonomy) of existing research on AI-based DDoS detection methods, together with an automatic clustering-based taxonomy, and compare the two.

3.1 DDoS attack classification

In Figure 1, we provide an overview of the four general categories of DDoS attacks, which we describe individually next. The most common attacks in real-world are volumetric DDoS attacks, in which an attacker coordinates bot networks to send a large number of UDP packets to a target victim. This increased traffic affects network and compute infrastructures, which now have to appropriately handle the increased fake traffic in detriment of traffic from legitimate users.

Next, protocol DDoS attacks exploit weaknesses in protocols themselves (mostly network layer protocols, such as TCP/IP/ICMP) and their software implementations and usually target the compute infrastructure of the victim operating system. This leads to an increase in the consumption of hardware resources to the point where the service becomes unable to handle legitimate requests.

Many of the attacks described in the volumetric and protocol categories can be further enhanced by reflection/amplification DDoS techniques. In this case, the attacker uses a bot network to take advantage of the connectionless nature of UDP to send requests with a spoofed IP address to multiple legitimate UDP-based services. The responses from these so-called reflectors, which are usually also amplified by a certain factor, overwhelm the target victim, further exacerbating the disruption of both network and compute capabilities.

Finally, application DDoS attacks target specific higher-level network services, including HTTP, web applications, and database management systems. In this scenario, an attacker sends a large number of well-crafted malicious requests to a network application that cannot properly handle the requests. This leads to increased consumption of hardware resources that ultimately make the application nonresponsive or outright crash the application, and thus denying the
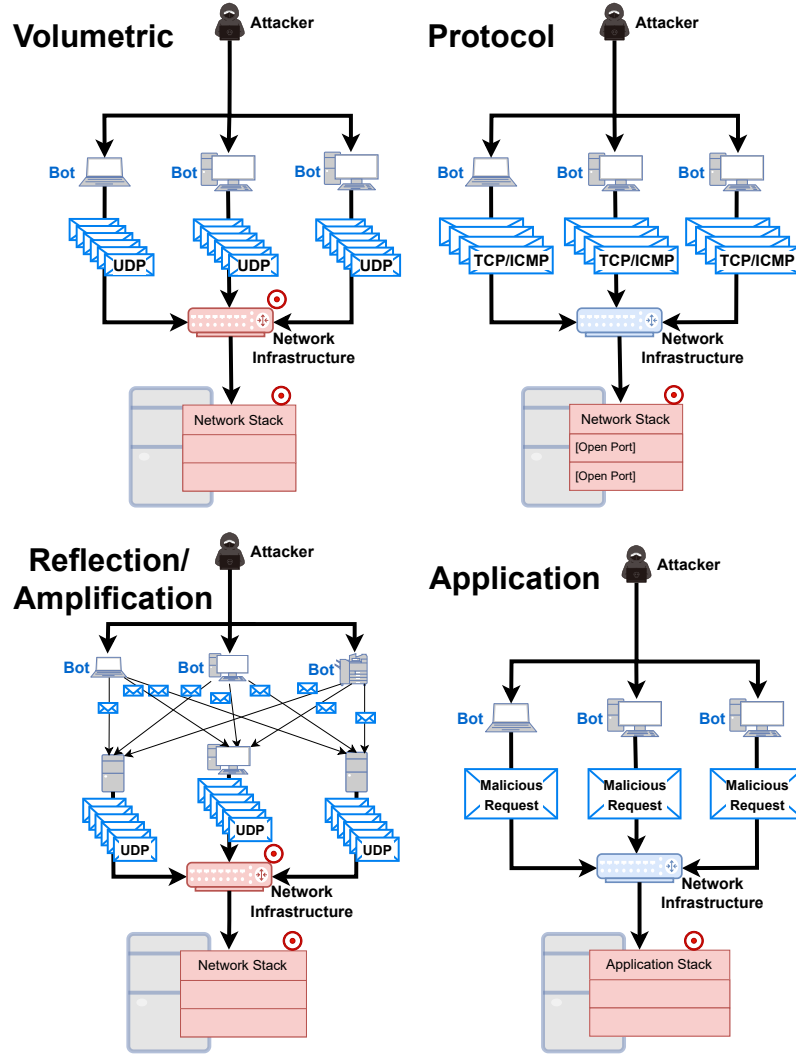
Fig. 1. Taxonomy of the four major categories of DDoS attacks described in separate diagrams. In each diagram, we depict an attacker who has control over networks of bot machines and targets a particular victim. Packets involved in the attacks bear the inscriptions of the typically used protocols, while on the victim side, we highlight in red and mark with bullseye the network resources that are the focus of the attacks: the network infrastructure such as communication hardware, the network software stack of the victim operating system, and the application software stack running user services.

services to legitimate users. In this scenario, the network infrastructure is mostly unaffected and the main target is the victim application stack on top of the operating system.

	Attack name	Target	Brief description
Volumetric	UDP flood (DNS, NTP, SSDP, IPsec, QOTD, SNMP, QUIC)	Network	Large number of UDP packets, sent via high-level protocols which are based on UDP
	ICMP fragmentation flood		Malformed ICMP packets
	Ping flood (IP/ICMP)		Large number of ICMP packets
Protocol	TCP flood (SYN/ACK/RST/FIN)	Network/ Compute	Large number of TCP (SYN/ACK/RST) packages, connections not fully established
	Ping of Death (IP/ICMP)	Network	Malformed ICMP packet(s)
	Smurf DDoS (ICMP flood)		Large number of spoofed ICMP packets
	BGP flood		Malformed Border Gateway Protocol packets, disrupt the routing tables
	DHCP Starvation (IP + MAC)		Large number of requests from spoofed MAC addresses which consume the DHCP IP pool
Reflection/Amplification	UDP flood (DNS, NTP, SSDP, QOTD, RPC, NetBIOS, CLDAP)	Network/ Compute	Large number of UDP packets, sent via high-level protocols which are based on UDP, redirected from multiple vulnerable machines or bot networks, manyfold amplified
	TCP flood (SYN/ACK/RST/FIN)		Large number of TCP (SYN/ACK/RST) packages, connections not fully established, redirected from multiple vulnerable machines or bot networks, manyfold amplified
Application	HTTP GET/POST flood (TCP flood)	Network/ Compute	Large number of concurrent GET/POST requests or send large files via HTTP
	HTTP Slowloris/R.U.D.Y.	Compute	Large number of concurrent slow or incomplete HTTP requests
	HTTP Initiated (HTTP Asynchronous)		Malicious HTTP requests that target web application logic/features
	DNS Server Query Flood		Large number of legitimate-looking DNS queries
	DNS NXDOMAIN/ Water torture		Large number of queries for non-existent domains
	ReDoS		Malicious regex search patterns that overload compute resources
	Database DDoS		Malicious HTTP requests that strain the DBMS

Table 3. A description of the most popular DDoS approaches from each of the four attack categories. We highlight the main attack target (network infrastructure and/or computational resources) and concisely describe the attack.

In Table 3, we provide a list of the most common DDoS attacks, grouped into each of the four main attack categories, together with a brief description and highlighting the main attack target: network infrastructure and/or computational resources.
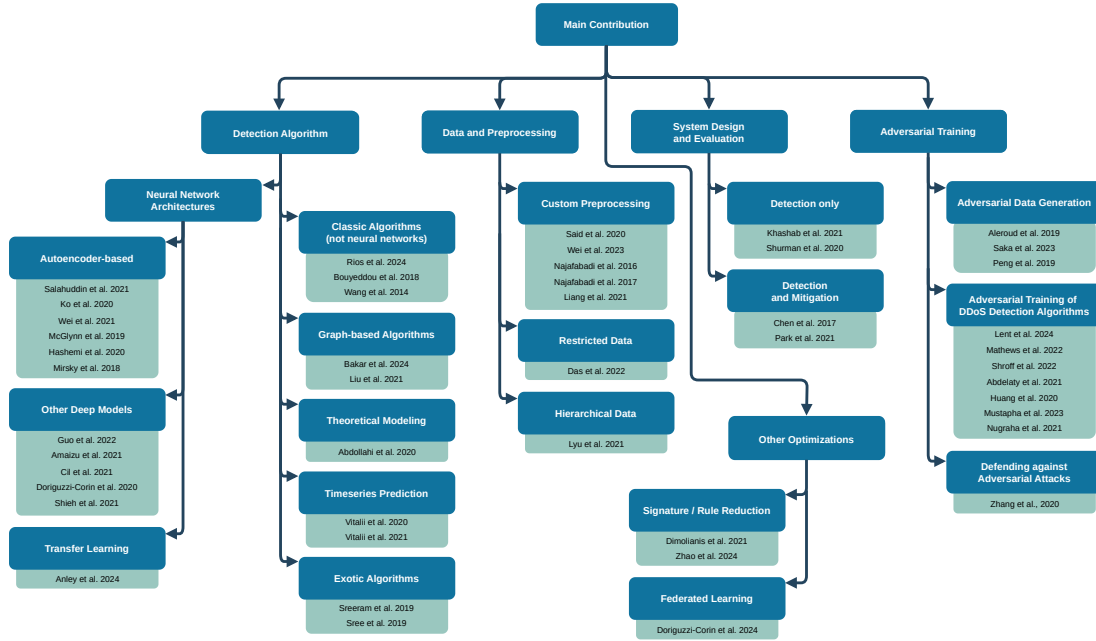
Fig. 2. Hierarchical grouping of the surveyed articles according to their main innovative contribution. References are linked to the papers (click on a reference to open its DOI page of the article).

3.2 Existing AI-based DDoS detection research

In an attempt to help the reader quickly navigate the state-of-the-art landscape and identify key affinities, we group the research papers according to their main innovative contribution, e.g., whether they focus on a particular detection algorithm, a new data preprocessing method, or building and evaluating a full DDoS mitigation system, etc. The resulting classification is presented in Figure 2, and contains five major groups. We acknowledge that research articles differ widely in breadth, scope, and innovative content, and it is sometimes difficult to assign them to a single category. The grouping simply means that we judged that particular aspect of the article as the most relevant in the overall context of the literature, which may help the reader grasp the main directions.

As illustrated in Figure 2, the largest group focuses on the actual algorithm for detecting DDoS attacks. The majority of the proposed approaches rely on deep neural networks, the autoencoder being a particularly popular architecture, thanks to its ability to detect deviations and abnormal patterns. However, several other custom architectures have been proposed, e.g., using bidirectional GRU layers with self-attention [35]. Of particular interest is study of Anley et al. [12], which also investigates how well the detection results transfer to other datasets than the ones used for training. Besides neural networks, various other algorithms are investigated, including classic supervised solutions (e.g., RFs), graph-based kernels [56], various modelings of traffic patterns, and even bio-inspired heuristic classification algorithms.

The second group contains papers with novelties related to data and preprocessing, ahead of the actual classification stage. A typical example of these custom methods is the approach of Wei et al. [114], which uses LSTM coupled with an autoencoder to extract time-varying features of the data, identifying anomalies using the autoencoder reconstruction

error. Another example is the method of Das et al. [29], which restricts the data used in the detection process only to port statistics, without any flow-related data, unlike the vast majority of related approaches.

Several articles, grouped under the "Other Optimizations" label, focus on collateral optimizations which may be highly relevant for the efficiency of DDoS solutions in practice. Thus, Dimolianis et al. [30] and Zhao et al. [125] consider the problem of consolidating attack signatures or mitigation rules, to detect and block several types of attacks simultaneously, with few rules. This also increases the robustness against new, unseen variants of the attacks. An adaptive federated learning approach is proposed in [32], facilitating cooperation between entities while avoiding the risk of exposing sensitive information.

The "System Design and Evaluation" group contains papers that, without focusing on a single particular aspect, are relevant to evaluations of DDoS detection pipelines. We differentiate here between articles focusing only on the detection, e.g. [47], and articles that also consider the mitigation process, e.g. [76].

The last group gathers articles involving adversarial data and training, as a way of increasing the robustness of anti-DDoS solutions. Since these articles typically use specific algorithms, rarely encountered in other articles (e.g., GAN architectures), we group all of them into a separate category. We further subgroup the articles according to whether they focus primarily on the actual data generation or the adversarial training process, with a separate note for the study of Zhang et al. [122], who investigate strategies against adversarial attacks (e.g., ensemble voting).

3.3 Manual vs. automatic taxonomies of AI-based DDoS detection methods

The dendrogram presented in Figure 3 is generated by applying an agglomerative clustering algorithm based on Ward's linkage over the TF-IDF representation of article titles and abstracts.

The TF-IDF representation is an extension of the simple bag-of-words model, which combines two components: term frequency (TF) and inverse document frequency (IDF). TF measures how often a word appears in a document, while IDF demotes common words by assigning lower importance to terms that appear in many documents. The result is a weighted vector representation that emphasizes unique and meaningful words for distinguishing documents.

Agglomerative clustering based on Ward's linkage is a hierarchical clustering algorithm that groups data points by iteratively merging clusters to minimize the total within-cluster variance. It starts with each TF-IDF vector as its own cluster and merges pairs of clusters, step by step. Ward's method chooses the pair of clusters to merge based on the smallest increase in variance within all clusters:

$$\Delta(C_i, C_j) = \text{Var}(C_i \cup C_j) - \text{Var}(C_i) - \text{Var}(C_j), \tag{1}$$

where C_i and C_j are two clusters, and the variance of a cluster is computed as follows:

$$\text{Var}(C_i) = \frac{1}{|C_i|} \sum_{x \in C_i} \Delta_{l_2}(x, \mu_{C_i}), \tag{2}$$

where x is a TF-IDF vector representing one of the surveyed articles, μ_{C_i} is the mean vector of cluster C_i, and Δ_{l_2} is the ED. Formally, the optimization criterion used to select the pair of clusters to be merged is defined as:

$$\arg \min_{i,j} \{\Delta(C_i, C_j)\}, \forall i \neq j. \tag{3}$$

This approach ensures that the resulting clusters are as compact and similar as possible, making it particularly effective for detecting spherical or evenly distributed clusters.

Fig. 3. A hierarchical clustering of the surveyed articles based on Ward's linkage. Each article is represented through a TF-IDF vector computed from the concatenated title and abstract of the respective article. The distance represented on the horizontal axis is computed via Eq. (1). The dendrogram is manually annotated to indicate meaningful groups of related papers. Best viewed in color.

Dataset	Attacks	Format	Raw Size (GB)	Flows Size (GB)
Bot-IoT [49]	Flood, TCP, HTTP	Raw, Flows	69.3	16.7
CAIDA 2007 [9]	Flood	Raw	21	-
CIC-DDoS2019 [93]	Flood, TCP, NTP, DNS	Raw, Flows, Timeseries	21.34	0.5
CICIDS 2017 [92]	HTTP, Slow	Raw, Flows	51.1	-
Darpa98 [54]	Flood, TCP, Smurf, HTTP	Raw	5.5	-
KDD99 [46]	Flood, Smurf, PoD	Flows	-	0.743
NSL-KDD [102]	TCP, Smurf	Flows	-	0.02
UNSW-NB15 [64]	Generic DoS	Raw, Flows	100	-
TON-IoT [8]	HTTP	Raw, Flows, Timeseries	25	-
TUIDS [13]	Flood, Smurf	Raw, Flows		private

Table 4. Available DDoS datasets with covered attack types, data size and data format. The attacks are sorted by attack class and then by attack type with volumetric first, protocol, and application last. Most datasets provide the raw format and optionally some preprocessed formats such as flows or timeseries; we also included the preprocessed data size where available. Dataset names link to the download site.

We manually labeled the resulting clusters after automatically obtaining the taxonomy, by investigating the grouped papers and finding the commonality from the point of view of a human expert. Next, we briefly walk from the outer clusters towards the inward labeled clusters and provide a few comments. We found that the first level splits between DNS, HTTP, and SYN flood attacks and also creates a clear split between Adversarial Training in normal and SDN contexts. Another fine line was the separation between Detection, a large cluster, and IDS focused research, a niche topic. Going deeper inside the Detection cluster, we find clusters separating shallow and deep learning, and also a cluster that grouped attack detection and attack classification papers. There is another shallow learning cluster within DNS Floods that was accurately separated from the general detection cluster. In the final clustering level, the manual labels provide a specialized view of the resulting dendrogram. For example, we have DNS Flood cache-based mitigation papers and adversarial training in SDNs employing active learning techniques.

It is interesting to compare these clustering results with the manually defined groups in Figure 2, in an attempt to discover common traits. A major similarity is that the "Detection" cluster in Figure 3 is almost completely included in the "Detection Algorithm" group in Figure 2, suggesting a clear and well-defined focus on detection methods. In addition, the two clusters of "Adversarial Training" in Figure 3 overlap greatly with that in Figure 2, which is to be expected given the specific algorithms and keywords used in this approach. Some other small similarities are visible as well, for specific niches, as in "timeseries forecasting/prediction".

Beyond these similarities, there are few shared insights. The main reason is that the criteria discovered in the automatic clustering process are diverse, including specific targeted protocols (HTTP, ICMP, DNS, SYN floods), generic approaches (Detection, Classification, Clustering), or network environments (e.g., IoT, SDN). On the contrary, in Figure 2, there is a single generic criterion considered, namely the technical approach proposed in each article. It is therefore natural that they lead to fairly different hierarchies. This, however, is not detrimental; it only helps the reader to obtain a more nuanced understanding of the surveyed literature.

4 Training Data: Flows, Graphs and Timeseries

We next discuss the various forms DDoS training data can take before being processed by learning algorithms. Most IDS analyze the traffic at various network layers. The first step in data processing covers the capture and storage of network traffic data. This step is essential for reaching high DDoS detection quality.

Existing public databases are gathered in Table 4. We can see that, except for KDD, all data are presented in the raw packet (pcap) format. Apart from CIC-DDoS2019, these datasets do not focus on DDoS attacks and instead include other types of attacks, besides normal traffic. Among DDoS traffic, flood attacks seem to be the most common, while TCP and HTTP come in second; each of these belong to the volumetric, protocol, and application attack classes, respectively. We also include, for brevity, the TUIDS database, which is not currently available to the public, although it was in the past and is often discussed in literature. Apart from raw traffic, some databases also provide preprocessed data in the form of flows or timeseries. We add this information in the table where available, and proceed to discuss these and other preprocessing formats.

In the network traffic structure, a *packet* is a data unit composed of *header* fields and the proper message (*payload*) to be transferred. Each field from the the *headers* stack, that belongs to a different TCP/IP layer, is necessary for the proper management and routing of the packet toward the destination. Capturing tools such as tcpdump are often used to monitor and store packet-level information of the network. When a capturing tool stores and inspects the traffic among multiple machines, the complexity of packet analysis is highly dependent on the number of monitored devices and on the common growing networking speeds. Additionally, the detection of real collective attacks (such as DDoS) usually requires a periodic inspection of very large low-level packet sets. For this reason, collection, inspection and management of low-level packet traffic remains a challenge for this format. Therefore, grouping packets into sets that have some shared characteristics becomes a natural idea.

A *flow* represents a packet collection related through a common set of attributes. The attributes include header fields corresponding to transport and application layers (IP addresses, IP protocol, source ports, etc.), as well as various information related to local packet management (e.g. input-output physical port). Despite the wide interval of attributes, collection and computation of flow records is an increasingly common capability of usual network devices. The flow records computed on the export device are transferred towards a collector that stores them. Each record has basic flow information and other fields: timestamp for flow start and stop, byte length over the packets in the flow, etc. This way, the flow format will eventually be much more compact than the packet data format.

Regardless of the primary traffic format, three dominant data models allow the diversity of anomaly detection techniques to be applied for detecting network intrusions: tabular, timeseries, and graph models.
Identifying important traffic attributes over each flow (or packets group) is a widely adopted preprocessing step that converts raw flow data into a new *tabular* dataset. After this conversion, many authors proceed to apply particular outlier detection techniques and provide specific interpretations of the results. Some common traffic attributes, retained from the raw traffic, are promoted by many papers to facilitate the prevention of multiple DDoS attacks. Thus, we further exemplify some of the most common features (that we found in the surveyed papers) organized over the class of DDoS attacks they address:

- Volumetric: source IP, destination IP, source port, protocol, source packet, destination packet, packet number (identifier), total number of bytes transferred, ToS values, duration of established connections, statistics from counting source MAC and IP address combinations, channel and socket properties.
- Protocol: source IP, source and destination IP combinations, IP identifier and fragmentation properties (DF field), TTL values, source and destination TCP port combinations, window size for TCP packets, maximum/minimum packet length, average packet size, maximum/minimum packet forward length.
- Application: incoming IP addresses, TCP source/destination port numbers, maximum number of sessions, web page access count, number of HTTP request packets per unit time.
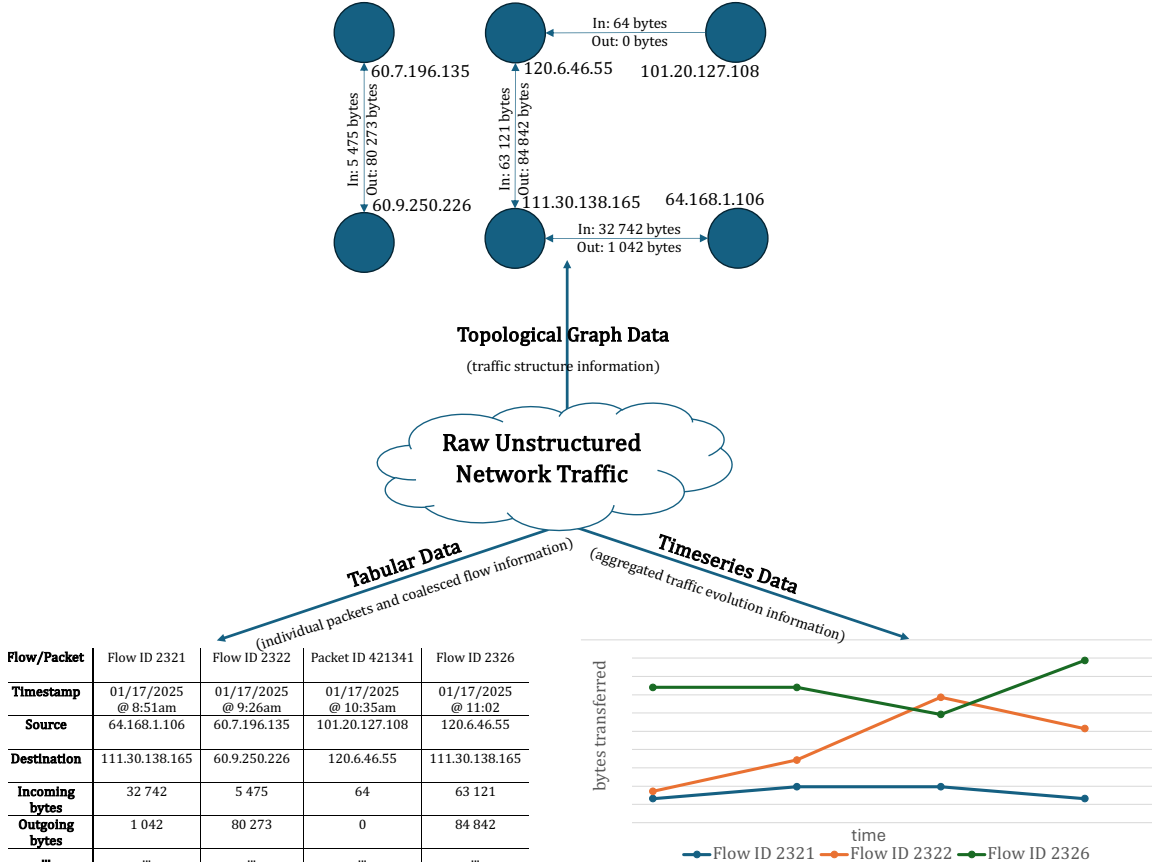
Fig. 4. Three representative toy examples of the most popular data extraction/organization techniques used for DDoS applications. From raw unstructured network traffic dumps, we can extract individual packet-level data, or coalesce the data into flows, and we can organize the data into: tabular, timeseries, and graph/network structured data.

By taking into account the timestamp of each flow, a temporal order can be established over the network traffic and the raw data takes the form of a *timeseries*. With a relevant encoding of the attributes, the DDoS attack detection problem becomes a timeseries analysis problem. Among the main flow attributes that prove to be worth tracking over time are: new flow record creations per second, average flow duration, average number of bytes per flow, average number of packets per flow [40]; DNS query rate [107]; statistics based on counting MAC and IP addresses combinations [63]; number of packets, entropy of IP addresses and ports list [52]; number of received packets, received bytes, number of sent packets, sent bytes, port keep alive duration, drop rate for received packets, number of errors for received and transmit packets, etc. [29].

The analysis of inherent relations between the flow attributes naturally leads to the third data model based on *graphs*. In general, DDoS attacks are prominently reflected/amplified in particular subgraph patterns where a single node has a high input degree compared with its neighborhood. Even a simple feature table of the flows, that includes a timestamp, source and destination IPs for each flow, can be easily converted into a basic static graph structure where the nodes are the IP addresses and the links are the connections within the network (see Figure 4).

5 Detection

We next delve into a detailed literature review of state-of-the-art AI-powered detection methods and solutions against DDoS attacks. Based on their specifics, we group attacks into volumetric, protocol-based, reflection and amplification, and application (Layer 7) attacks. We also summarize the reported detection time of the proposed methods from the papers that include this information.

5.1 Volumetric DDoS attacks

Volumetric attacks are conectionless floods in which an attacker uses networks of bots to send a large number of, usually, UDP packets towards an intended victim, disrupting mainly their network infrastructure, but potentially also their network stacks.

In the vast majority of cases, volumetric attacks are treated in a unified manner, just as the other DDoS attacks. This is because well-established datasets, such as CIC-DDoS2019 [93], have recorded traffic from various types of DDoS attacks. As stated in [22], given the CIC-DDoS2019 dataset, there are two main approaches: (1) detect that a DDoS attack is happening and (2) establish which kind of DDoS attack it is. In general, when discussing synthetic datasets, such as CIC-DDoS2019, the results published in the literature are excellent, exceeding 99% precision[2] for detecting DDoS attacks using only relatively simple AI algorithms. Most probably, these almost perfect results stem from the synthetic nature of the employed datasets, since these were generated in synthetic lab conditions, not real-world situations.

Most research papers obtain similar near-perfect accuracy results. For example, Wei et al. [113] use a mixed architecture composed of an autoencoder and a multi-layered perceptron. In this case, the autoencoder performs dimensionality reduction, while the perceptron performs the actual classification task.

A few-shot model is one that aims to recognize attacks after only seeing very few examples during training. If there are no training examples at all, the model is called zero-shot. This scenario is very useful in the real world because new types of attacks (or variations of the existing ones) appear all the time. Because of that, an important characteristic of an AI model is its ability to detect and respond to new types of attacks. For instance, Shieh et al. [94] propose a human-in-the-loop, semi-supervised solution. New types of attacks are detected in a zero-shot manner. The relevant data samples are sent for validation and manual classification to a human operator. If what is detected is indeed a new type of attack, the information is added to the system, such that, in the future, the AI system will be able to automatically perform the correct classification for this type of attack. The implementation is based on a combination of deep learning, BiLSTM, and GMM techniques.

In most solutions, traffic is modeled using network flow models, making explicit use of the structure of the network. For example, in [56], the traffic is modeled through a time-varying network graph model. The regular network traffic is modeled as a fixed set of normalized graphs. An extra element of novelty is that, to measure the divergence from the regular traffic, the authors use a Weisfeiler-Lehman kernel. Compared with previous solutions, the proposed AI system performs better in real-time scenarios and for the classification of various DDoS attacks. In another graph-based approach [3], a two-stage aggregation technique is proposed, resulting in the construction of graph models for the packet-level and the flow-level traffic. The authors analyze this traffic using CNNs, and obtain state-of-the-art results (especially in terms of the F1 score). The accuracy of the model is proportional to the past traffic window size, a parameter that is important in establishing the topology of the traffic information, but whose increase leads to significantly larger inference times.

---

[2]All other metrics, such as accuracy, recall or F1 have similarly high values over 99%.

Detection of DDoS attacks plays a crucial role in many practical situations. For example, smart camera surveillance systems are critical because disabling them via a DDoS attack allows an on-the-ground operation to proceed undetected. Hence, the network robustness of such systems is essential. The DDoS robustness of surveillance systems is analyzed in the study of Mirsky et al. [63], which is based on a set of serially connected autoencoders. Their role is to encode the traffic characteristics and prepare them for anomaly detection methods. For training, the authors use custom datasets encapsulating a set of relevant cyberattacks including recon, man-in-the-middle, and DDoS (SYN flood, SSDP flood, and SSL renegotiation). The proposed solution is highly applicable and focuses on many practical considerations.

5.2 Protocol DDoS attacks

Protocol-based DDoS attacks generally exploit specific characteristics of network protocols or their software implementation (e.g., TCP, ICMP). In general, the main target of these attacks is the victim's network stack, but the network hardware infrastructure might also be significantly affected.

Among the earliest work in protocol based DDoS detection using ML techniques includes the detection of SIP flood attacks over VoIP networks. First, the work in [71] uses an SVM approach and 38 features grouped into four categories to detect telephony flooding and spam attacks. The features are constructed from time-sliced statistics of the signaling messages usually used in SIP for VoIP (especially INVITE messages). Then, the work in [4] uses ML techniques such as NB and DT on a dataset composed of both spatial and temporal engineered features extracted from raw network packets. In this setting, temporal features are variations in the number of INVITE, ACK, and BYE messages captured in traffic while spatial features are computed by the entropy of caller IDs in a time window. The results reported in both papers are above 90% in terms of detection accuracy.

As mentioned in previous sections, many solutions in the scientific literature jointly handle many DDoS attacks, developing general detection and protection methods working with multiple attacks. This is also reflected in the databases used for training and validation [92, 93]. More precisely, even if individual attacks differ, technically depending on the specifics of the protocol exploited, the types of extracted features and the used detection methods are usually generic, and can be applied to a wide range of attacks. For example, in [125], the datasets used for evaluation contain 89 different variations of DDoS attacks. Even though some works explicitly refer only to certain types of attacks, usually TCP or TCP-SYN floods [29, 30, 35], this is largely due to the test environment and the datasets used for evaluation, and not to the methods themselves being specialized for certain types of attacks. In contrast, specializing on certain types of attacks can be inherently problematic, since several types of attacks are not represented well enough in the current databases to be properly modeled and analyzed [29].

Most of the works analyzed in this section start from features extracted from flows, such as the number of packets, number of bytes, and duration [47], together with identification elements (source address, source port, destination address, destination port, and protocol). The values are taken at periodic intervals to monitor their evolution in time. Several works mention the particular relevance of these differential parameters: Khashab et al. [47] use the number of flows to the same host and the same port, respectively, from the last 5 seconds; Das et al. [29] take the differences between the number of packets and bytes at successive intervals of 5 seconds; Zhao et al. [125] use a set of 5 sliding time windows, ranging from 0.1 to 10 seconds. The papers above use manually defined features, but other approaches use automatic methods for extracting relevant features, such as those based on autoencoders [48], feed-forward networks [53] or attention mechanisms [35].

A special case is the work of Das et al. [29], which uses only parameters collected from the ports in the network infrastructure, without considering features extracted from flows. This is considered an advantage, as flow-level statistics

are dependent on the network topology, and the associated traffic, which could impact the generalization of the results [29]. In this case, once an attack is detected, the origin of the attack is identified through a localization process, which analyzes traffic on all ports and identifies the switches most affected, considered to be the closest to the source of the attack.

The methods used for detection are varied, but most of them are based on supervised learning algorithms: SVM, RF, DT, neural networks with MLP or CNN architectures [27, 29, 47, 125]. Along the same lines, the work in [103] uses a simple OC-SVM classifier to detect particular types of DHCP starvation attacks. The approach is rather simple and uses four features extracted from DHCP traffic: the number of DISCOVER, REQUEST, and DECLINE DHCP messages and the total number of ARP messages in a fixed time quanta. These statistics are enough to reach near-perfect detection accuracy. The majority of works use and compare several types of classifiers, with sometimes extremely different results, without a consensus on the best method (for example, Das et al. [29] report an accuracy of 51% with MLP networks). An exception to this is the articles in which the classification method is linked with the feature extraction method, such as [48], which uses autoencoders for feature extraction simultaneously with classification. Most classification methods are traditional feature-based approaches, in different flavors. An exception is the approach of Guo et al. [35], which uses bidirectional GRU models, coupled with a hierarchical attention mechanism.

Many of the works that use supervised methods report detection accuracy values of over 99%, often over 99.9%, for example [29–31, 35, 125]. It is not clear to what extent these performance levels transfer to real test environments, or are just a consequence of using the limited set of publicly available databases, or of using certain software tools for simulating attacks in tests. However, we note that many works mention evaluating the results on real, private datasets recorded from ISPs [30, 125], which increases the level of confidence in the proposed solutions.

Besides attack detection, a large number of works also consider the other stages required in the entire mitigation pipeline, proposing various optimizations and improvements. A popular topic is optimizing the number of rules required for traffic filtering, which may increase proportionally with the number of addresses involved in the attack, potentially becoming very costly [30, 125]. Several approaches to consolidate the filtering rules are proposed, which rely on the similarity of some attack patterns to counter them with as few rules as possible. In [30], a method based on genetic algorithms is proposed, which can reduce the number of rules and signatures used in detection by up to 99.99%, according to the authors. We note, however, that this work considers only SYN Flood attacks. These consolidated rules are implemented with programmable XDP enabled firewalls; since optimizing the rules is a slower process, the SYN cookies technique is used as a temporary protection measure for quick response, until the rules are updated.

The concept of "DDoS attack family" is proposed in [125], aggregating attacks that allow similar defense strategies. This is done by evaluating the differences between signatures of different attacks, representing the relationships between them as graphs, followed by partitioning them into sub-communities with specific algorithms. The approach also allows for adaptation to new types of attacks, not encountered in the design phase. The analysis of the authors shows that 89 types of attacks can be grouped into just four families, namely: i) attacks to network Layer 4 or higher; ii) TCP attacks without a response from the victim; iii) TCP attacks with a response from the victim, and; iv) connectionless attacks. As proof of the efficiency of this solution, to counter all attacks in category iii), only two rules are needed, which are based on counting RST packets and packets having both fields SYN and ACK set.

The possibility of adapting to new types of attacks is also addressed in the work of Doriguzzi-Corin et al. [32], who investigate a federated learning approach to incorporate data and models from multiple clients, while avoiding data centralization. In [53], a semi-supervised learning method is introduced, which is appropriate in situations where

attacks are poorly labeled or even unknown. The method is based on clustering the data and then projecting it into a reduced space, aiming to preserve the cluster membership for data with available labels.

As mentioned above, even though most methods do not specialize on a specific attack, some do specialize. In the following, we address these specialized methods. For TCP Connection/SYN Flood, in addition to the articles presented in the previous section, those using the CIC-DDoS2019 dataset [93] implicitly address SYN Flood. Thus, the works of Salahuddin et al. [88], Wei et al. [113], and Shieh et al. [94], described in Section 5.1, are relevant to the subject. PoD attacks are characterized by packets having, after reassembly, a size greater than the maximum size of 65535 allowed by the protocol. This attack is explicitly addressed in [2], which detects it using a threshold for the number of packets in a time window with a size greater than a certain limit. Next, the BGP attack exploits the vulnerabilities of the BGP protocol, and is based on sending malicious routing updates. In [60], these attacks are detected using two autoencoders to extract the essential features from the AS-specific information. Finally, Smurf DDoS attacks involve spoofing the source IP addresses of ICMP packets sent to broadcast addresses, which causes a large number of responses to be sent to the devices holding the spoofed IP addresses. Nowadays, this type of attack is simply avoided by disabling the broadcasting option on network devices. For this reason, there are few works addressing this attack in the literature, and most of them use relatively simple approaches, such as [17], which harnesses the Kullback-Leibler divergence to identify anomalies. Instead, when it comes to the ICMP protocol, interest in Ping Floods have resurfaced recently in the context of IoT networks [7].

5.3 Reflection and amplification DDoS attacks

RA-DDoS attacks enhance the previously described volumetric and protocol DDoS attacks. Historically, reflection and amplification attacks represented a subset of the broad class of volumetric attacks and are characterized by an attacker who can selectively use a relatively small amount of traffic employing bots, vulnerable protocols and services to generate a considerable amount of traffic toward an intended victim. These types of attacks usually target services such as DNS and NTP, as well as both standard transport protocols: mostly UDP, but also TCP. An attacker could just reflect traffic toward a victim, but by far, the most dangerous attacks are those where the attack volume is amplified manyfold. Although RA-DDoS attacks were initially based exclusively on UDP, subsequent advances [14, 51] showed how TCP can also be used to reflect and amplify malicious traffic. Due to these characteristics, reflection and amplification attacks are sometimes treated separately in the DDoS volumetric attack literature [62].

Most classic techniques used to detect and mitigate RA-DDoS attacks rely on changing the network's configuration parameters or the targeted services. This mitigation technique bears the name System and Network Hardening [14, 50, 51, 84]. Anticipating AI-based techniques, a particular line of work [80, 107, 109] follows the idea of collecting simple statistics about network traffic to detect and then suggest mechanisms to stop RA-DDoS attacks.

Among the first papers to use AI techniques to detect RA-DDoS attacks, we mention here works [19, 61] that use several machine learning techniques such as DT, MLP, NB, and SVM to detect DNS RA-DDoS attacks.

In the following decade, Anley et al. [12] turned to modern AI techniques, such as deep CNNs together with transfer learning, to enhance the generalization capabilities of AI models by combining data from multiple freely available DDoS datasets. The recent literature views the problem of generating realistic RA-DDoS attacks as a central one in the development of reliable detection methods. For this reason, the work of Mathews et al. [59] proposes a robust RA-DDoS method that uses neural networks and adversarial learning techniques, such as EAD and TextAttack, to generate counterexamples for the AI model, which will not be correctly labeled as attacks. Such enhancements lead to better generalization in practical datasets. In the same spirit, a method that uses a GAN-based discriminator for

anomaly detection is proposed in [52]. This uses GRU-type neurons, analyzing traffic as a time series. Features such as the number of transmitted packets, the number of bits and entropy of IP addresses and ports are extracted for each one-second time window, and are aggregated into overlapping 10-second time series for analysis. The generator network learns to synthesize normal traffic flows with features very similar to those from the training set, while the discriminator improves the anomaly detection capability. In addition to the detection module, the authors also design a mitigation module that identifies a victim as the IP address that receives the largest number of flows in a time window detected as an anomaly, and subsequently blocks the IP addresses that send traffic towards it.

Regarding the features used in the detection process, several approaches are possible. One of them involves selecting a subset of features through different methods, as in [96, 114]. Another option, found in [52, 58], is to build new features from a time window containing a variable number of packets or flows. The former option allows for inference as soon as the information about a flow is stored in the database, whereas the latter requires building new features only after all the flows corresponding to a time window are available and stored. Amaizu et al. [11] use the Pearson correlation coefficient for feature selection and then jointly train two neural networks with different architectures to predict the type of attack for the flows in the CIC-DDoS2019 dataset. The attack classes are not limited to RA-DDoS types, as they also include general volumetric attacks.

As with other types of attacks, RA-DDoS attacks are usually addressed collectively in the scientific literature. A notable exception is the work of Lyu et al. [58], who develop detection methods for RA-DDoS attacks based on the DNS protocol. They propose a hierarchical graph structure with a dynamic retention policy to model traffic at three different levels: host, subnet, and AS. Then, two different methodologies are used to detect traffic anomalies at each level. The first method uses static detection rules for each scenario, while the second uses IF and OC-SVM models to detect anomalies based on the following time-windowed features: variance of packet size, number of internal hosts queried by each external entity, average number of packets sent in queries to each internal host, and variance of the number of packets sent in queries to them. The authors further used two of these features to separate anomalous hosts in scanners and flooders.

Because of the time-dependent nature of monitoring network traffic quantities, recurrent neural structures have been used with great success for RA-DDoS detection. Shurman et al. [96] proposed two approaches to detect RA-DDoS attacks. The first one is a simple framework consisting of a database with known attack signatures and an anomaly-based detector. Here, the signature database is updated when a packet whose characteristics deviate from the normal ones is analyzed by the anomaly-based detector. The second approach uses LSTM-based networks with varying number of LSTM layers to detect RA-DDoS attacks from the CIC-DDoS2019 dataset. An RF classifier is used together with the GINI impurity to select a subset of features for training the LSTM-based model. Another LSTM-based architecture is proposed in [114] to detect RA-DDoS attacks as time series anomalies within a time window containing a configurable number of data flows. The training process uses a subset of the features from legitimate data flows, and then a threshold on the reconstruction errors obtained by the LSTM autoencoder is derived. The proposed architecture contains an encoder, in which the output of each LSTM cell component is ignored, followed by a decoder, where the outputs of the component cells represent the reconstructions of the input data. The authors reported separate results for RA-DDoS attacks based on DNS, LDAP and SNMP. A similar LSTM autoencoder is used in [86], but here, latent representations are used to train an OC-SVM model, because the simple threshold-based approach did not achieve satisfactory performance. In both papers, only instances labeled as normal are used in the training process. While the work of Wei et al. [114] was used to detect three types of RA-DDoS attacks, Said et al. [86] presented a model trained to detect several types of attacks, including DDoS.

A recent survey [44] focuses particularly on RA-DDoS attacks, and yet another one [108] showcases the recent work in studying RA-DDoS attacks in the context of SDNs.

5.4   Application-level DDoS attacks

We provide an overview of the current state of research as well as some possible future research directions concerning application-level DDoS attacks. Alongside the information present in the surveys section and the works concerned with Layer 7 Application attacks, we can distinguish three broad types of attacks, each of them having multiple subtypes. More precisely, we discuss detecting HTTP attacks, especially flooding attacks, followed by detecting slow HTTP attacks (Slowloris and R.U.D.Y.). Lastly, we discuss hierarchical DNS server attacks.

As compared with previous DDoS attacks, application layer DDoS usually targets the operating system, software stack and hardware compute resources of the target victim machines. The network infrastructure is under some strain, but it is usually not the main target of the attack. In many application layer DDoS attacks, the issue is not necessarily the large number of packets transferred, but the compute burden due to the malicious content of the messages and how these are interpreted and processed by the victim's software.

5.4.1   HTTP/2, HTTP/HTTPS flood. HTTP flood attacks refer to the scenario where an attacker uses malicious HTTP requests to target an HTTP server with the goal of making it unresponsive to legitimate traffic.

Because the HTTP requests involved in this class of attacks are perfectly normal requests, those attacks are difficult to tell apart from the legitimate user traffic. Wang et al. [110] propose a method to classify HTTP flood attacks into two types: random and perfect knowledge. In perfect knowledge ones, the attackers know the structure of the website and attempt to mimic the behavior of a legitimate user. For the random ones, this assumption does not apply. This difference is important for the proposed solution, because it uses the learned distribution of the requests to identify the legit users. The solution, named HTTP-SoLDiER, is based on the correspondence between the individual clicks generated by users with the general distribution of the users of the site. Based on deviation theory, the solution can differentiate legitimate users from attackers by computing the probability of navigation, with respect to general interest and popularity. The critical point of the paper is the fact that, because the content of the pages and the structure of the site can change, the probability must be dynamically adjusted. To prevent the corruption of statistics during an attack, the authors propose a correction algorithm. This algorithm, called EWMA, estimates the popularity of pages in real-time, based on current state and historical data, thereby reducing the bias introduced by the attack itself. A noteworthy observation, also mentioned in previous literature, is that the general distribution of page popularity follows a Zipf law [79]. Because the complexity of their proposed solution is computationally $O(M)$, where M is the number of pages, the computation requirements are not significant. Alongside the correction algorithm, the success rate in detecting malicious traffic is 99% for random HTTP floods and 77. 9% for perfect-knowledge attacks. The downside of this solution is that some websites only have a single page/URL and the so-called circular perfect-knowledge attacks, where the attackers mimic the behavior of a legitimate user by rotating several pages, are significantly harder to detect.

Anomaly detection can be used to detect HTTP flood attacks. For example, Najafabadi et al. [68] propose an anomaly detection mechanism, based on text mining, to detect HTTP flood attacks. To be close to a real-world scenario, the benign data is generated by collecting the internal traffic in the university network and the attack traffic is generated during a penetration testing session designed to simulate a real attack. Despite that, the dataset is still a synthetic one, generated using software tools. The proposed mechanism works as follows: a document is defined as a request-response series corresponding to an HTTP GET. The requested resources are regarded as words/tokens. Features are extracted

from each document by using bi-grams. OC-SVM is trained using benign data, to model the normal behavior of users. Then, this model is tested on attack data to measure performance.

Inspired by bioinformatics, some algorithms for the detection of HTTP floods are tested by Sreeram et al. [43] and Sree et al. [100]. The study of Sreeram et al. [43] is a theoretical one, using the CAIDA dataset. The employed metrics include session time, the maximum number of sessions, the minimal time between requests, and the number of different packet types observed during this session window. The metaphor-based metaheuristics Bat algorithm is used for detection in both papers. The algorithm simulates bats that fly in a search-space, and the distance and direction of the search are dynamically corrected based on the intensity and frequency of the echolocation pulses.

Sree et al. [100] propose using the same Bat algorithm as a clustering method to find sources of HTTP flood attacks. This paper focuses mainly on the cloud computing aspects, detailing the practical side of the implementation. The method used in the paper, fuzzy Bat clustering, combines FCM and the meta-heuristic algorithm inspired by the Bat algorithm. This combination is designed to tackle the problems of FCM, such as the selection of the cluster center, detecting unknown attacks, and avoiding local optima. The proposed software solution consists of four components: log collection, data storage, processing, and the detection method itself. Data are collected from virtual machine logs, network logs, and access logs. Once stored, the data is cleaned by pre-processing, and fed to the detection module. The parameters used for classification and identification are: the number of requests in a time window, the frequency of the requests, response times, request similarity, and the number of times the same request is repeated. The testing framework is private, in the cloud, using OpenStack. The malicious traffic is generated using software tools.

DDoS HTTP mitigation methods are grouped by Park et al. [76] into two categories: destination level mitigation (implemented by the HTTP server) and network level mitigation. The work proposes an architecture designed to mitigate both DDoS HTTP flood attacks. The solution uses a combination of SDN and rules implemented on the web server. This way, by using both mechanisms, after initially identifying an attacker at the server level, the subsequent requests can be interrupted promptly, before reaching the web server, reducing the impact of the attack. The paper presents the general architecture without providing details about the implementation, the detection algorithm, or about the restrictions imposed on the infrastructure during the mitigation process. The hardware solution consists of two OpenFlow switches connected to each other, one placed between the attacker and the server and the other one close to the server. The switches are connected to an SDN controller. The controller is named HDFD and it has multiple roles in mitigating the attack damage. The protection flow operates like this: once a request is classified as suspicious by the server-level detection component, this connection is sent to the SDN. The switches interrupt the server connection and the SDN mimics the server, responding using headers similar to those of the server, to trick the attacker. Those headers are useful to restore the connection if the user is in fact legitimate. If the user does not manage to pass the test that was set up by the SDN, then the rules of the switches are updated, blocking the suspicious connection.

5.4.2 Low & slow attacks: R.U.D.Y. and Slowloris. In HTTP flood attacks, the sheer number and size of requests are used to overload the server. R.U.D.Y. and Slowloris attacks are more sophisticated because they use specially crafted connections. More precisely, Slowloris transmits the data very slowly, and R.U.D.Y. fails to complete the post requests. The similarity is that they both mimic clients with slow Internet connections, which can clog the computational resources of the victim's web server.

The performance of a few basic ML algorithms in detecting R.U.D.Y. attacks is compared by Najafabadi et al. [69]. This paper offers a scalable solution and compares three different detection algorithms: k-NN, C4.4D and C4.5N DTs. The attribute selection is decided by comparing 10 different methods, such as: F-metrics, geometric means, mutual

information, Fisher scores, Kolmogorov-Smirnov statistics, Chi Squared, S2N, among others. Tests are performed on traffic taken from SANTA, a dataset composed of commercial traffic collected in the network of an ISP. It is worth mentioning that the collected traffic is bidirectional. By using Netflows, the author can group the data into sessions corresponding to complete round-trips. The selected parameters are divided into three categories: session similarity, periodicity, and the speed corresponding to requests and responses. To compare the predictive models, the author uses ANOVA analysis, thereby measuring the performance with the full feature set versus the characteristics selected in the comparison process.

Another method to detect the low & slow DDoS HTTP attacks is by using time series, as shown by Vitalii et al. [89]. The idea of this work is to use time series to predict normal user behavior. Using these statistical predictions, a projected trajectory of the future actions performed by the user is computed. Those computations are based on a set of selected parameters. This mechanism is often found in literature, based on four components: a module for collecting traffic, data storage, pre-processing of data and traffic parameter computation and, finally, a detection module. The authors emphasize the importance of detecting attacks early. They also highlight the impossibility of completely preventing attacks, due to the fact that they mimic normal user behavior to a large extent. It is worth mentioning that, because slow attacks are the opposite of flood attacks, the parameters used for computations are different from the ones used in the previous section. Predicting the latency of the network packets enables the detection of slow DDoS attacks, based on an algorithm that identifies the unknown future values in a traffic parameter time series. The method mixes self-learning and statistical analysis, based on the existence of a sufficient volume of statistical data about the attacks.

A solid state-of-the-art analysis is given in Rios et al. [83]. The work also proposes a new way of detecting slow HTTP DDoS attacks, especially Slowloris. Because Slowloris traces are publicly unavailable, the authors collected data in three different media: emulated, LAN, and MAN. To facilitate the 10-fold cross-validation, the authors combine four data traces (emulated, LAN, MAN-IF and MAN-Pal) in a single dataset. This dataset is randomly split 75%/25% for each of the 10 folds. The original solution is a combination of three methods: FL, RF and ED. Because of this, the proposed method is called FRE. FL is initially used for classification, RF for a detailed analysis, and ED is a fail-safe mechanism to break the ties. If the first two methods share the same result, it is considered final, otherwise ED is used to compute the minimal distance between features, comparing the normal and attack patterns. Nine ML algorithms are analyzed in the paper using two scenarios, one with generic hyper-parameters and the other with optimized hyper-parameters. The analyzed algorithms are: GNB, MNB, XGB, LGBM, k-NN, MLP, SVM, DT, and RF.

5.4.3 DNS server query attacks. This type of DDoS attack targets the DNS server infrastructure of a network. It can target root servers, top-level-domain servers, name-servers, and sometimes even resolvers. Such attacks are hard to detect and stop because, often, these servers are not under the control of the administrators and developers of the services that are the ultimate targets of the attack.

There is a common frustration caused by the difficulty of any major upgrade of the DNS protocol and the slow, difficult or missing coordination among the DNS providers to implement the proposed DNS infrastructure and solutions [112]. This is why, in time, public institutions and private enterprises have migrated towards implementing private DNS servers to facilitate testing and data gathering, and to benefit from the flexibility of a custom implementation. Most solutions recognize the difficulty in changing or updating the DNS protocol and focus instead on backward-compatible optimizations that can be implemented on a subset of servers to diminish the currently known issues [101].

Fig. 5. Papers that reported DDoS detection time (left) and their associated accuracy score on testing data where available (right). The papers are sorted in descending order with citation indexes on the bottom axis, author names on the bar (including DOI links) and publication year at the top. The bars are color coded: blue for deep learning and green for shallow learning methods. The red line shows top reported accuracy results.

As the technology advanced and the adoption of AI increased, the direction of research in DNS flood prevention also changed. It is worth mentioning, though, that the works proposing AI-based solutions start from the same premise, more precisely, the control over the DNS protocol, either in the resolver part or at another point in the chain.

The work of Lyu et al. [58], also covered in the RA-DDoS section, proposes a new method to detect DDoS DNS attacks by implementing a hierarchical graph structure. The solution can classify, with a high degree of certainty, the following attacks: scans, DNS floods, and slow attacks. This paper presumes control over the DNS resolvers and of the top level domains. The solution is implemented on the network boundary, by duplicating the traffic using an SDN switch. The traffic is subsequently parsed and only the DNS-related one is kept. A pre-processing module is implemented using the DPDK and Intel NFF-Go libraries. The parameter values are combined and estimated to cover the whole range of attacks. The new data is added to a dynamic graph to be analyzed. The resulting graph has a hierarchical structure that is able to establish a link between the attack targets and the IP addresses, IP classes, and the zones where the attacks originate. Hence, the anomaly detection decisions can be stratified, leading to better overall precision.

5.5 Detection delays

The damage caused by a DDoS attack is highly determined by the necessary duration of a given system to detect and mitigate the attack. Thus, in addition to accuracy or other learning quality indicators, a particularly important feature of detection systems is the detection time (or speed) of an incoming attack. Even though this aspect is quite insightful, we found a relatively small amount of papers that rigorously report the detection time or speed of their particular

methods or systems. Referring to detection time, most authors report some values for their particular context (dataset dimension, hardware configuration, methodologies, etc.) that are measured in seconds, and in rare cases, in-flows per second or samples per second.

Few papers report the delay between mounting the DDoS attack and its detection. In Figure 5, we plot a comparison of the papers that do. We tried to keep a fair comparison where we take the best results presented by the authors when performing detection on testing data, including the raw packet processing time. Some authors do not discuss or include this last bit of information. Where available, we also added the associated accuracy for the reported detection time. From the publication years at the top, we can observe that recent years show an improvement in detection times, even though this is not a strong trend. It is also visible that shallow learning methods dominate. Regarding the accuracy-time trade-off, we see very high accuracy values, which can lead to simple selection criteria: choose the method with the fastest detection time. It will be interesting to see how the accuracy quality fares when using the trained model across different databases, which we discuss as a future direction at the end of this survey.

6   AI Generated DDoS Traffic

Training robust systems for DDoS attack detection involves the use of a rich dataset of examples with a high degree of variety. Since obtaining representative data is not always easy, a few methods in recent literature have resorted to the use of generative models to obtain synthetic training data, with the objective of replacing the tedious process of collecting real data. Among the employed models are GANs and VAEs.

Some methods have employed autoencoders as the main tool for improving network IDS by handling noisy or incomplete data. For example, Hashemi et al. [37] propose a method to improve DDoS detection systems using autoencoder-based architectures to eliminate the noise (denoising autoencoders). The authors studied two architectures, RePo (Reconstruction from Partial Observation) and RePo+. The first one aims to reduce the adversarial traffic reconstruction error by training the model for traffic reconstruction from incomplete data, which leads to a better classification between benign and malicious attacks. RePo+ uses stochastic inference with multiple random masks to reduce the chances of adversarial examples fooling the system. The result of using denoising autoencoders for improving network IDS is the increase of detection accuracy by up to 45% in the adversarial context. Overall, the detection of attacks improved by 29% on the packet level and 10% on the stream level, compared with the results obtained by methods such as Kitsune, DAGMM, BiGAN. In a more recent yet similar approach, Saka et al. [87] extend the use of autoencoder architectures and create entirely new traffic samples with the TVAE, specifically designed for generating synthetic tabular data. As opposed to standard VAEs, which struggle with handling mixed data types, TVAE learns a probabilistic latent space representation of the data, ensuring that the newly generated samples maintain the original dataset's structure and dependencies. Despite the success in mirroring the real data distribution, the synthetic dataset used to train a RF Classifier registered a performance drop in accuracy to 93%, compared with the 98%-99% achieved by GAN-based models.

While autoencoder-based methods focus on improving the quality of generated data, GANs offer an alternative approach by generating realistic data which enhances model robustness. In Figure 6, we illustrate a typical GAN-based approach for improving the efficiency of intrusion detection models. Dual approaches, such as the ones in [87], bridge the methods by also highlighting the strengths of GAN variants, specifically CTGAN and CopulaGAN. CTGAN uses conditional input during training to model feature dependencies, while CopulaGAN employs preprocessing with Gaussian copula transformations.
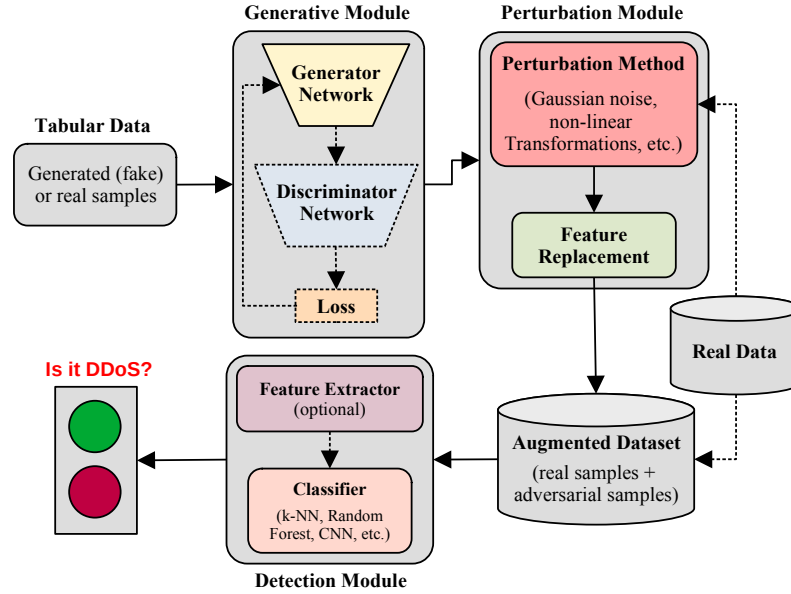
Fig. 6. A generic GAN-based architecture for a network intrusion detection system. The input typically consists of tabular data, which can be either real samples from the dataset or generated fake samples. This data is sent through the Generative Module, which comprises a GAN designed to create realistic synthetic data that mimics benign traffic or DDoS attacks. This module can also contain preprocessing steps, such as ternary encoding [6] or conditional information [87]. The role of the Feature Perturbation Module is to alter the previously generated samples (creating malicious examples that look benign or slightly perturbed versions of the original attacks) using Gaussian noise, non-linear transformations, or other methods. The output of this process is an augmented dataset, which combines the original real samples and the newly perturbed adversarial samples. The augmented dataset is further used to train and/or evaluate a detection system. In the training scenario, a classifier is used to measure the robustness against the adversarial examples.

An earlier application of GANs, SDN-GAN [6], represents a specialized method designed to generate synthetic attacks targeting software-defined networks. Unlike traditional GANs, this variant employs a binary encoding for representing features. Moreover, it goes beyond the classic GAN framework (based on a generator and a discriminator) by incorporating a new component, the intrusion detector, which creates a feedback loop that produces adversarial samples that can bypass detection. Without GAN-generated adversarial examples, the initial experiments using traditional models achieved relatively high detection rates (81%-85%). Consequently, detection rates dropped to as low as 7% (with a maximum accuracy of 53% for RF) when adversarial examples were introduced. While the findings are relevant for understanding vulnerabilities in software-defined networks, the used dataset (CAIDA 2007 DDoS attack dataset) includes older traffic samples that might not resemble newer environments and techniques, limiting comparability with recent work featuring modern datasets. Abdelaty et al. [1] tackle this problem by evaluating their model on the CICIDS2017 dataset, a more comprehensive and diverse benchmark. Their study is also focused on generating high-quality samples and showcasing the impact of using generative models for generating network traffic against machine learning models used by tools for DoS/DDoS attack detection. The proposed architecture is based on two components. The first component is a generative model trained using normal traffic captures (WGAN-GP), aimed to generate additional traffic that respects the normal distribution of the features extracted from the training set. The role of the second component is to replace the features of normal traffic with features from the generation process.

The authors show that including a GAN in the training stage of a model used to detect attacks leads to an F1 score improvement of 0.32, compared with the results obtained by the reference models.

6.1 Adversarial training for DDoS detection

Adversarial training is a neural network training technique that aims to enhance the robustness of neural networks by exposing them to adversarially perturbed data in the training phase. One way of achieving this is through dynamic adversarial training, which involves combining a primary objective function which needs to be minimized, with a secondary objective function (adversarial), which ought to be maximized. This training approach is meant to avoid learning certain examples or features that can affect the generalization of the model if they end up being learned by the model. Adversarial training requires careful application because it can affect the convergence of the model to an optimal point. For this reason, the optimization steps for minimizing the primary objective function are generally larger than the steps to maximize the secondary objective function. This can be achieved by weighting the objective functions so that the weight of the primary function is larger than that of the secondary function. Another way to achieve the same effect is by using different learning rates. Minimizing the primary objective function can be accomplished by applying the gradient descent algorithm. In theory, the secondary objective function is maximized by the gradient climbing/ascent. In practice, this behavior can be simulated by still applying the gradient descent algorithm, but changing either the sign of the objective function, the learning rate, or the gradients.

A practical example of dynamic adversarial training can be seen in the work of Zhang et al. [122], who propose a deep learning-based defensive mechanism for IDS, called Tiki-Taka. In the first part of the paper, the authors test the vulnerabilities of three existing IDS models that use neural networks (MLP, CNN, and C-LSTM) against five recent black-box adversarial attacks. The result shows that up to 35.7% of attacks successfully bypass the models, as evaluated on the CSE-CIC-IDS2018 dataset. The proposed improvement of the method targets three defensive mechanisms: model voting ensembling, ensemble adversarial training, and adversarial query detection. The first mechanism combines the performance of multiple classifiers to reduce the probability of an attack passing through the system unnoticed. The second one aims to continuously augment the dataset with adversarial examples and retraining the models. Finally, query detection adds another layer of security by blocking the attacker's IP address in case an attack is detected. The authors show that these methods are more efficient in their evaluations, with an increase in detection accuracy to almost 100%.

In contrast, static adversarial training, as implemented by Nugraha et al. [73], involves pre-generating adversarial examples (using mechanisms such as the FGSM and JSMA) before training and incorporating them into the dataset. The paper proposes a deep learning method to detect DDoS attacks, trained both with and without adversarial data. The authors employed two architectures, namely CNN-LSTM and MLP. When exposed to adversarial validation data, the MLP and CNN-LSTM models recorded drops of 11.99% and 9.87%, respectively. To address these performance drops, the authors attempted to train the MLP (chosen because of its efficiency) using three adversarial training procedures. Out of the tested methods, the one that involves replacing 80,000 examples from the training set with adversarial data achieves similar results to the originally obtained accuracy (99%).

6.2 Adversarial examples and attacks

In the current literature, some methods for detecting and classifying DoS and DDoS attacks use different machine learning algorithms, including deep learning. With the exponential increase in traffic volume, algorithms such as SVMs, DTs, and Bayesian networks became inefficient for training and testing. Current directions in the architecture

of network IDS favor deep learning methods because of their capacity to precisely classify traffic in a network, by learning certain abstract representations from a large volume of data. However, since many of these models are trained on a single dataset, training sets and subsets originate from the same source. Thus, if the input data comes from an external source, with some small change in features, algorithms may fail to generalize. Since most machine learning algorithms are vulnerable to adversarial examples, lately, we have noticed attackers using generative AI techniques to generate them and produce incorrect classification decisions in standard detection systems. Even if the properties of an intrusion detection system are not known (black box), potential attackers can generate adversarial examples by repeatedly changing small subsets of features in the traffic (e.g., the time interval between consecutive packets). After each iteration, based on the received answer (confirmation/blocking of the attack or lack of an answer), the perturbations can be adjusted or changed completely until the network is breached. In this way, malicious fluxes are classified as benign traffic and remain undetected by the model.

The vulnerabilities of DoS IDS based on artificial neural networks against black-box adversarial attacks became an extensively studied matter. Peng et al. [77] propose an attack method that employs generating adversarial DoS data by disrupting relevant traffic features, achieving a synthetic data generation success rate of 80.77%. Their experiments on KDDCup99 and CICIDS2017 datasets indicate high initial classification accuracies of 98.69% and 93.45%, respectively. These values drop significantly in the presence of adversarial examples. While Peng et al. [77] highlight the vulnerabilities of detection systems using synthetic adversarial data, Huang et al. [41] extend the approach by proposing two methods for generating DDoS adversarial data, which can both be directly applied to LSTM-based models: GA and PWPSA. The GA technique uses a genetic algorithm to gradually evolve derived data from the original data into DDoS adversarial attacks, with success rates between 72.50% and 99.09%. PWPSA employs an interactive process of modifying the data, determining the position and packet that affect the classification the most (based on a predefined function), with success rates between 67.17% and 95.35%. Overall, GA achieves better results and is also preferred in practice, while PWPSA is more computationally efficient, with a slightly lower success rate than GA.

To address the challenges created by adversarial traffic, recent studies have turned to GANs for both generating adversarial data and improving network IDS. Mustapha et al. [67] showcases a method to detect DDoS attacks based on LSTM neural networks, which achieves 100% accuracy on an initial dataset. This approach is further tested against adversarial attacks generated with a GAN model, where a significant decrease in performance is noticed. The final solution builds on the initial architecture by creating two new models: the first detects adversarial generated traffic, while the second classifies the traffic as normal or DDoS. The simultaneous use of the two techniques led to a final detection accuracy of 91.75%. On a similar note, Shroff et al. [95] initially use WGAN-GP (as Abdelaty et al. [1]) to generate benign traffic captures, as well as DDoS attack captures, which are then used to test the accuracy of some classifiers. Furthermore, the data is used to experiment with a 5-layer deep network, achieving an accuracy of 95.37% in identifying DDoS attacks generated by GANs. The authors show that manually modifying DDoS attack-specific features in benign data leads to an improved training method in the context of adversarial attack detection.

He et al. [38] find that feature-level attacks such as FGSM and JSMA do not generate practical adversarial data. Furthermore, dependence on outdated datasets, like SL-KDD and KDDCup99, decreases the relevance of results. More recent datasets, such as CICIDS2017/2018 provide better benchmarks, but indicate the need for further development in adversarial traffic generation. White-box attacks are highly effective in evading detection systems, but black-box attacks still struggle because of the increased traffic complexity. As also noted by Abdelaty et al. [1] and Saka et al. [87], GAN-based architectures show potential in adversarial data generation. Nevertheless, He et al. [38] conclude that the

literature still needs more representative and diverse datasets, along with more robust defense mechanisms such as adversarial training and feature reduction strategies.

7   AI Generated Mitigations

Once a DDoS attack is mounted, we are continuously pressured to act and to act fast. After detecting the attack, we quickly have to mitigate it through various techniques, most of which involve blocking the attack nodes while allowing legitimate traffic to navigate as usual. The complexity of existing attacks and network topologies requires us to write sophisticated and efficient blocking rules: we want to block as many connections as possible within a single firewall rule such that throughput is not throttled. The last couple of years have shown an emergent trend that passes the responsibility of compiling block lists and composing efficient filtering rules from cybersecurity agents to specialized AI-agents that build custom-made DTs or fine-tune LLMs to generate firewall rules.

Louro et al. [57] establish a baseline by employing existing LLMs, both commercial and open-source, to tackle the mitigation task, a task at which they show limited to no success rates. Here, mitigation consists of emitting proper `iptables` and `snort` rules for various DDoS attacks that are passed as (structured) prompts. Significant improvements are shown by employing LoRA of LLMs and PEFT techniques to fine-tune existing open-source models, Mistral and LLaMA, for this mitigation task, which result in outputs consisting of advanced filtering rules with high success rates. Following the same process, ShieldGPT [111] provides an AI-based DDoS mitigation software architecture that classifies incoming DDoS attacks and fine-tunes GPT for prompt templating `iptables` rules as required by each attack type.

Some approaches, such as DrLLM [119], avoid fine-tuning by employing prompt engineering techniques like CoD, for templating the output, and Zero-shot CoT, for processing incoming flows as online timeseries. The authors prove the success of this approach on popular vanilla models like GPT, LLaMA, Qwen2 and Deepseek. While DrLLM handles solely DDoS detection, regardless of the attack type, this approach could be further extended to the mitigation tasks discussed above.

If so far we have seen flows based on tools for tools, LLMs that produce rules for packet filters, there is also literature on fine-tuning LLMs for human agents under DDoS attack. Indeed, in ShieldGPT [111], the authors design a a separate sub-task to provide an explanation prompt template for the cybersecurity agent handling the case. The work of Păduraru et al. [74] introduces CyberGuardian based on a fine-tuned LLaMA model in order to aid agents in tackling various cybersecurity attack scenarios. In the DDoS scenario, 23 agents participate in simulated scenarios where they prompt the LLM to achieve two tasks: generating IP block lists and emitting appropriate firewall commands to stop the attack.

A very popular ML technique to learn AI generated mitigations to DDoS attacks is DTs. These models are very popular in our setting because they are highly interpretable, their results can relatively easily be transformed into logical rules based on *and/or* operations, and they can be trained very fast as compared with other ML techniques from the literature. The work in [121] proposes a DT model to infer filtering rules to mitigate volumetric DDoS attacks. The authors explain how to convert the DT model into interpretable logical rules and they obtain state-of-the-art results. The work in [23] proposes a method called Anomaly2Sign which is based on DT models and generates Suricata rules for a wide range of DDoS attacks. The authors also focus on training the simplest DT models, measured by AIC scores, in order to further improve the interpretability and simplicity of the learned rules. To train their DT models, both papers assume that the available dataset is labeled and separated (albeit not exactly perfectly) into legitimate and illegitimate traffic data. In addition, both papers highlight the low training times, below one second, for their DT models and tout their computational efficiency compared with other ML models. This is a crucial factor when considering real-time or online training scenarios. Finally, both papers report near-perfect mitigation results for their proposed methodologies.

8 Conclusions and Future Research

In this survey, we provided a comprehensive review of AI-based detection and mitigation techniques for volumetric, protocol, application, reflection and amplification attacks. The study introduced a clear definition of these attacks and offered a manual and automatic taxonomy of existing research. Available datasets and data formats were discussed together with alternative formats such as timeseries, graphs and tabular data formats which are meant to aid the learning process. Of special note are the sections that treated AI generated attack mitigations and AI generated traffic for adversarial training and adversarial examples and attacks. This work leads to multiple insights into future research venues that we describe below.

Cross-dataset testing. A remaining open question concerns how well is the detection quality maintained across new, potentially unseen datasets. The risk of literature overfitting is valid, given the relatively small number of popular datasets used in most papers, and although private datasets are sometimes reported in the literature, the sensitive nature of network traffic and the difficulty of capturing relevant attacks are significant impediments. A recent investigation [12] reports drops in detection accuracy in the range of 5% - 10% for different models, when evaluating on datasets other than the ones used for training. Solutions based on supervised learning are, admittedly, affected to a greater extent compared to unsupervised approaches. Yet, the issue is largely under-addressed in literature, especially considering the impact it may have when deploying anti-DDoS systems in the wild.

Anti-DDoS tailored algorithms. Our literature review has shown limited specialization or data adaptation regarding AI algorithms. Most of the work was focused on handling cybersecurity attacks in bulk, as they were found in the various available datasets and the few works that specialized in DDoS attacks used standard shallow or deep learning algorithms for the task. We consider that future directions can improve results in terms of accuracy and, more importantly, detection and mitigation times, if new anti-DDoS tailored algorithms are investigated. Improvements will probably be more visible when coupled with cross-dataset testing.

Dynamic data format adaptation. DDoS attack bandwidths can vary widely. While most are centered around 1Gbps, recently we often find 5-10Gbps reports in the wild. A common attack behavior is that DDoS start small and grow toward their peak, as more bots are enabled, reflection and amplification are enabled, etc. Current literature does not cover how AI algorithms should handle this ramp-up and how traffic data storage and processing impacts detection and mitigation quality. For instance, one can easily imagine that storing raw packet data or even flows when under 10Gbps DDoS attacks is an almost impossible task: storing each packet (even in memory) for AI inference purposes can lead to self-DoS-ing the in-memory or on-disk database. Instead, algorithms need to be able to adapt and handle multiple data formats depending on the attack size: for example we can keep raw data until the attack reaches 1Gbps, switch to flows afterward until 2Gbps, and then shrink the data from IP-level to IP-class level until 5Gbps, and finally switch to a handful of statistics that can be quickly computed, stored and inferred.

Detection and mitigation time for algorithm quality assessment. Catching DDoS attacks early on is critical to keeping the victim infrastructure running. Still, almost the entire literature tackles this task as a standard classification problem: choose a dataset, train a classifier, test and validate the results and report standard metrics such as accuracy, F1-score, area under the curve, true positive rate, true negative rate, etc. While this is indeed necessary for basic model validation, we consider that the domain-specific validation should be detection time coupled with mitigation time and effectiveness, where applicable. More on mitigation, we consider that besides mitigation time, metrics regarding normal infrastructure

operation and regular traffic throughput while under DDoS should be measured and reported by future research when proposing new mitigation algorithms and methodologies.

Fuzzy detection. Most solutions use AI to partition the traffic into two distinct classes: normal and attack. In contrast, the prediction of ML algorithms is in fact a probability score, and the binary class label is determined via a probability threshold. Based on this observation, the detection threshold can be dynamically adjusted in order to maximize service availability. For example, when a server load is low, one can decrease the threshold such that more traffic is passed through and false positives are reduced. Subsequently, during high load the threshold can be increased. From another perspective, this can be seen as a scheduling algorithm where the classification score acts as traffic priority.

Comprehensive DDoS dataset. Most public datasets contain network recorded DDoS traffic covering a small subset of attack types, often mixed with non-DDoS attacks. We consider that there is a need for dedicated datasets that encompass all known DDoS attack types as well as more unusual benign traffic, like games, various messaging services, peer-to-peer exchanges, blockchain and SSH, to name a few. The traffic should be labeled and include further annotations regarding current throughput, bots details, reflection and amplification. The datasets should include various bandwidth traffic shapes, such as ramping up (e.g. from 100Mbps to 10Gbps), maintaining a fixed bandwidth (e.g. 5Gbps for 30 minutes), sine behavior (e.g. up and down from 1Gbps to 3Gbps for 2 hours). Expert mitigation rules that can be employed at various times to handle the attack should also be provided in order to evaluate mitigation algorithms and the firewall rules they provide.

Furthermore, we identify the need to have more complex and realistic datasets that mimic the real-world practical difficulties of detecting DDoS attacks. According to classic AI performance metrics, many current AI-based detection methods perform nearly perfectly on the available datasets. We believe that, at least partially, these results also betray the simplistic nature of the available datasets. Finally, training detection methods on current datasets do not generalize well to new DDoS attacks, even when these are variations on classic DDoS attacks. It is therefore desirable to have an adversarial approach to dataset design such that we do not overfit detection methods for particular scenarios and we thus prepare for real-world scenarios.

Adversarial training. Adversarial training is another area that is insufficiently explored. We consider that adversarial training techniques could broaden the applicability and robustness of AI-based DDoS attack detection methods to out-of-distribution samples. For example, in computer vision, adversarial domain adaptation [105] was shown to boost results in cross-domain settings. We consider that such an approach is also likely to perform well in DDoS attack detection for cross-dataset scenarios.

Explainable AI.. An understudied topic in the area of DDoS attack detection is the development of explainable AI methods. As in other domains, the success of deep learning methods comes with an important downside, namely that the reasons behind the predictions are often unknown or hard to precisely determine. The literature in this direction is scarce [10, 28], lacking sufficient exploration towards explainable DDoS attack detection based on various types of deep neural networks. Knowing the reasons behind a decision could also be useful in the mitigation phase.

AI mitigation. AI generated firewall rules are still an early research topic. While existing work has shown that it is a valid pursuit, much work can be done in this direction. Currently, algorithm efficiency is estimated by comparing the generated outputs with expert written rules for specific one-time attacks. We consider that this can be further improved by continuously generating and updating firewall rules, while maintaining the context of a prolonged attack. During

the attack, the algorithm has to take into consideration updating, merging or removing existing rules, not just inserting new ones (e.g. we want to block an entire IP class instead of keeping 255 separate block rules). The main goal here is to keep the firewall chain rules efficient, such that throughput is optimized; the victim infrastructure has to be kept running and regular traffic has to pass through as close to the network's nominal functionality as possible.

References

[1] Maged Abdelaty, Sandra Scott-Hayward, Roberto Doriguzzi-Corin, and Domenico Siracusa. 2021. GADoT: Gan-based adversarial training for robust DDoS attack detection. In *2021 IEEE Conference on Communications and Network Security (CNS)*. IEEE, 119–127.

[2] Asrin Abdollahi and Mohammad Fathi. 2020. An intrusion detection system on ping of death attacks in IoT networks. *Wireless Personal Communications* 112, 4 (2020), 2057–2070.

[3] Rana Abu Bakar, Lorenzo De Marinis, Filippo Cugini, and Francesco Paolucci. 2024. FTG-Net-E: A hierarchical ensemble graph neural network for DDoS attack detection. *Computer Networks* 250 (2024), 110508. https://doi.org/10.1016/j.comnet.2024.110508

[4] Muhammad Akbar and Muddassar Farooq. 2014. Securing SIP-based VoIP infrastructure against flooding attacks and Spam Over IP Telephony. *Knowledge and Information Systems* 38 (02 2014). https://doi.org/10.1007/s10115-012-0595-5

[5] Huda Ali Alatwi and Charles Morisset. 2021. Adversarial machine learning in network intrusion detection domain: A systematic review. *arXiv e-prints* (2021), arXiv–2112.

[6] Ahmed AlEroud and George Karabatis. 2020. SDN-GAN: Generative Adversarial Deep NNs for Synthesizing Cyber Attacks on Software Defined Networks. In *On the Move to Meaningful Internet Systems: OTM 2019 Workshops: Confederated International Workshops*. Springer, 211–220.

[7] Omar Mohammed Almorabea, Tariq Jamil Saifullah Khanzada, Muhammad Ahtisham Aslam, Fatheah Ahmad Hendi, and Ahmad Mohammed Almorabea. 2023. IoT Network-Based Intrusion Detection Framework: A Solution to Process Ping Floods Originating From Embedded Devices. *IEEE Access* 11 (2023), 119118–119145.

[8] Abdullah Alsaedi, Nour Moustafa, Zahir Tari, Abdun Mahmood, and Adnan Anwar. 2020. TON_IoT telemetry dataset: A new generation dataset of IoT and IIoT for data-driven intrusion detection systems. *IEEE Access* 8 (2020), 165130–165150.

[9] Sabah Alzahrani and Liang Hong. 2018. Generation of DDoS attack dataset for effective IDS development and evaluation. *Journal of Information Security* 9, 4 (2018), 225–241.

[10] Ahmad Alzu'bi, Amjad Albashayreh, Abdelrahman Abuarqoub, and Mai A. M. Alfawair. 2024. Explainable AI-Based DDoS Attacks Classification Using Deep Transfer Learning. *Computers, Materials and Continua* 80, 3 (2024), 3785–3802. https://doi.org/10.32604/cmc.2024.052599

[11] Gabriel Chukwunonso Amaizu, Cosmas Ifeanyi Nwakanma, Sanjay Bhardwaj, Jae-Min Lee, and Dong-Seong Kim. 2021. Composite and efficient DDoS attack detection framework for B5G networks. *Computer Networks* 188 (2021), 107871.

[12] Mulualem Bitew Anley, Angelo Genovese, Davide Agostinello, and Vincenzo Piuri. 2024. Robust DDoS attack detection with adaptive transfer learning. *Computers & Security* 144 (2024), 103962.

[13] Monowar H Bhuyan, Dhruba K Bhattacharyya, and Jugal K Kalita. 2015. Towards Generating Real-life Datasets for Network Intrusion Detection. *Int. J. Netw. Secur.* 17, 6 (2015), 683–701.

[14] Kevin Bock, Abdulrahman Alaraj, Yair Fax, Kyle Hurley, Eric Wustrow, and Dave Levin. 2021. Weaponizing middleboxes for TCP reflected amplification. In *30th USENIX Security Symposium (USENIX Security 21)*. 3345–3361.

[15] Nizar Bouguila and Wentao Fan. 2020. *Mixture models and applications*. Vol. 530. Springer.

[16] Anis Bourou, Valérie Mezger, and Auguste Genovesio. 2024. GANs Conditioning Methods: A Survey. *arXiv preprint arXiv:2408.15640* (2024).

[17] Benamar Bouyeddou, Fouzi Harrou, Ying Sun, and Benamar Kadri. 2018. Detection of smurf flooding attacks using Kullback-Leibler-based scheme. In *2018 4th International Conference on Computer and Technology Applications (ICCTA)*. IEEE, 11–15.

[18] Hamparsum Bozdogan. 2000. Akaike's information criterion and recent developments in information complexity. *Journal of mathematical psychology* 44, 1 (2000), 62–91.

[19] Chih-Chieh Chen, Yi-Ren Chen, Wei-Chih Lu, Shi-Chun Tsai, and Ming-Chuan Yang. 2017. Detecting amplification attacks with Software Defined Networking. In *2017 IEEE Conference on Dependable and Secure Computing*. 195–201. https://doi.org/10.1109/DESEC.2017.8073807

[20] Pin-Yu Chen, Yash Sharma, Huan Zhang, Jinfeng Yi, and Cho-Jui Hsieh. 2018. Ead: elastic-net attacks to deep neural networks via adversarial examples. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 32.

[21] Junyoung Chung, Caglar Gulcehre, Kyunghyun Cho, and Yoshua Bengio. 2014. Empirical evaluation of gated recurrent neural networks on sequence modeling. In *NIPS 2014 Workshop on Deep Learning, December 2014*.

[22] Abdullah Emir Cil, Kazim Yildiz, and Ali Buldu. 2021. Detection of DDoS attacks with feed forward based deep neural network model. *Expert Systems with Applications* 169 (2021), 114520.

[23] Antonio Coscia, Vincenzo Dentamaro, Stefano Galantucci, Antonio Maci, and Giuseppe Pirlo. 2024. Automatic decision tree-based NIDPS ruleset generation for DoS/DDoS attacks. *Journal of Information Security and Applications* 82 (2024), 103736.

[24] Vinícius G Costa and Carlos E Pedreira. 2023. Recent advances in decision trees: An updated survey. *Artificial Intelligence Review* 56, 5 (2023), 4765–4800.

[25] Padraig Cunningham and Sarah Jane Delany. 2021. K-nearest neighbour classifiers-a tutorial. *ACM computing surveys (CSUR)* 54, 6 (2021), 1–25.

[26] Federal Bureau of Investigation Cybersecurity and Infrastructure Security Agency. 2024. Understanding and Responding to Distributed Denial of Service Attacks. https://www.cisa.gov/resources-tools/resources/understanding-and-responding-distributed-denial-service-attacks

[27] Mohamed Amine Daoud, Youcef Dahmani, Mebarek Bendaoud, Abdelkader Ouared, and Hasan Ahmed. 2023. Convolutional neural network-based high-precision and speed detection system on CIDDS-001. *Data & Knowledge Engineering* 144 (2023), 102130.

[28] Saikat Das, Namita Agarwal, and Sajjan Shiva. 2021. DDoS Explainer using Interpretable Machine Learning. In *2021 IEEE 12th Annual Information Technology, Electronics and Mobile Communication Conference (IEMCON)*. 1–7. https://doi.org/10.1109/IEMCON53756.2021.9623251

[29] Tapadhir Das, Osama Abu Hamdan, Shamik Sengupta, and Engin Arslan. 2022. Flood Control: TCP-SYN Flood Detection for Software-Defined Networks using OpenFlow Port Statistics. In *2022 IEEE International Conference on Cyber Security and Resilience (CSR)*. 1–8. https://doi.org/10.1109/CSR54599.2022.9850339

[30] Marinos Dimolianis, Adam Pavlidis, and Vasilis Maglaris. 2021. SYN Flood Attack Detection and Mitigation using Machine Learning Traffic Classification and Programmable Data Plane Filtering. In *2021 24th Conference on Innovation in Clouds, Internet and Networks and Workshops (ICIN)*. 126–133. https://doi.org/10.1109/ICIN51074.2021.9385540 ISSN: 2472-8144.

[31] R. Doriguzzi-Corin, S. Millar, S. Scott-Hayward, J. Martínez-del Rincón, and D. Siracusa. 2020. Lucid: A Practical, Lightweight Deep Learning Solution for DDoS Attack Detection. *IEEE Transactions on Network and Service Management* 17, 2 (June 2020), 876–889. https://doi.org/10.1109/TNSM.2020.2971776

[32] Roberto Doriguzzi-Corin and Domenico Siracusa. 2024. FLAD: Adaptive Federated Learning for DDoS attack detection. *Computers & Security* 137 (Feb. 2024), 103597. https://doi.org/10.1016/j.cose.2023.103597

[33] Junliang Fan, Xin Ma, Lifeng Wu, Fucang Zhang, Xiang Yu, and Wenzhi Zeng. 2019. Light Gradient Boosting Machine: An efficient soft computing model for estimating daily reference evapotranspiration with local and external meteorological data. *Agricultural water management* 225 (2019), 105758.

[34] Benyamin Ghojogh, Ali Ghodsi, Fakhri Karray, and Mark Crowley. 2021. Factor analysis, probabilistic principal component analysis, variational inference, and variational autoencoder: Tutorial and survey. *arXiv preprint arXiv:2101.00734* (2021).

[35] Xiaojun Guo and Xuan Gao. 2022. A SYN Flood Attack Detection Method Based on Hierarchical Multihead Self-Attention Mechanism. *Security and Communication Networks* 2022, 1 (2022), 8515836. https://doi.org/10.1155/2022/8515836

[36] Zeyu Han, Chao Gao, Jinyang Liu, Jeff Zhang, and Sai Qian Zhang. 2024. Parameter-efficient fine-tuning for large models: A comprehensive survey. *arXiv preprint arXiv:2403.14608* (2024).

[37] Mohammad J Hashemi and Eric Keller. 2020. Enhancing robustness against adversarial examples in network intrusion detection systems. In *2020 IEEE Conference on Network Function Virtualization and Software Defined Networks (NFV-SDN)*. IEEE, 37–43.

[38] Ke He, Dan Dongseong Kim, and Muhammad Rizwan Asghar. 2023. Adversarial machine learning for network intrusion detection systems: A comprehensive survey. *IEEE Communications Surveys & Tutorials* 25, 1 (2023), 538–566. https://doi.org/10.1109/COMST.2022.3233793

[39] Paul Francis Hitesh Ballani. 2008. Mitigating DNS DoS Attacks. *CCS '08: Proceedings of the 15th ACM Conference on Computer and Communications Security* (2008), 189–198.

[40] Rick Hofstede, Václav Bartoš, Anna Sperotto, and Aiko Pras. 2013. Towards real-time intrusion detection for NetFlow and IPFIX. In *Proceedings of the 9th International Conference on Network and Service Management (CNSM 2013)*. IEEE, 227–234.

[41] Weiqing Huang, Xiao Peng, Zhixin Shi, and Yuru Ma. 2020. Adversarial attack against LSTM-based DDoS intrusion detection system. In *2020 IEEE 32nd International Conference on Tools with Artificial Intelligence (ICTAI)*. IEEE, 686–693.

[42] Rami Ibrahim and M. Omair Shafiq. 2023. Explainable Convolutional Neural Networks: A Taxonomy, Review, and Future Directions. *ACM Comput. Surv.* 55, 10, Article 206 (Feb. 2023), 37 pages. https://doi.org/10.1145/3563691

[43] Venkata Praveen Kumar Vuppala Indraneel Sreeram. 2019. HTTP flood attack detection in application layer using machine learning metrics and bio inspired bat algorithm. *Applied Computing and Informatics* 15, 1 (2019), 59–66.

[44] Salih Ismail, Hani Ragab Hassen, Mike Just, and Hind Zantout. 2021. A review of amplification-based distributed denial of service attacks and their mitigation. *Computers & Security* 109 (2021), 102380.

[45] Mohamed Riadh Kadri, Abdelkrim Abdelli, Jalel Ben Othman, and Lynda Mokdad. 2024. Survey and classification of Dos and DDos attack detection and validation approaches for IoT environments. *Internet of Things* 25 (April 2024), 101021. https://doi.org/10.1016/j.iot.2023.101021

[46] H Günes Kayacik, A Nur Zincir-Heywood, and Malcolm I Heywood. 2005. Selecting features for intrusion detection: A feature relevance analysis on KDD 99 intrusion detection datasets. In *Proceedings of the Third Annual Conference on Privacy, Security and Trust*, Vol. 94. Citeseer, 1723–1722.

[47] Fatima Khashab, Joanna Moubarak, Antoine Feghali, and Carole Bassil. 2021. DDoS Attack Detection and Mitigation in SDN using Machine Learning. In *2021 IEEE 7th International Conference on Network Softwarization (NetSoft)*. 395–401. https://doi.org/10.1109/NetSoft51509.2021.9492558 ISSN: 2693-9789.

[48] Ili Ko, Desmond Chambers, and Enda Barrett. 2020. Adaptable feature-selecting and threshold-moving complete autoencoder for DDoS flood attack mitigation. *Journal of Information Security and Applications* 55 (Dec. 2020), 102647. https://doi.org/10.1016/j.jisa.2020.102647

[49] Nickolaos Koroniotis, Nour Moustafa, Elena Sitnikova, and Benjamin Turnbull. 2019. Towards the development of realistic botnet dataset in the internet of things for network forensic analytics: Bot-IoT dataset. *Future Generation Computer Systems* 100 (2019), 779–796.

[50] Marc Kührer, Thomas Hupperich, Christian Rossow, and Thorsten Holz. 2014. Exit from Hell? Reducing the Impact of Amplification DDoS Attacks. In *23rd USENIX security symposium (USENIX security 14)*. 111–125.

[51] Marc Kührer, Thomas Hupperich, Christian Rossow, and Thorsten Holz. 2014. Hell of a handshake: abusing TCP for reflective amplification DDoS attacks. In *8th USENIX Workshop on Offensive Technologies (WOOT 14)*. 1–6.

[52] Daniel M Brandão Lent, Vitor G Da Silva Ruffo, Luiz F Carvalho, Jaime Lloret, Joel JPC Rodrigues, and Mario Lemes Proença. 2024. An Unsupervised Generative Adversarial Network System to Detect DDoS Attacks in SDN. *IEEE Access* (2024).

[53] Junjie Liang, Wenbo Guo, Tongbo Luo, Vasant Honavar, Gang Wang, and Xinyu Xing. 2021. FARE: Enabling Fine-grained Attack Categorization under Low-quality Labeled Data. In *The Network and Distributed System Security Symposium 2021*. https://www.ndss-symposium.org/ndss-paper/fare-enabling-fine-grained-attack-categorization-under-low-quality-labeled-data/

[54] Richard Lippmann, Joshua W Haines, David J Fried, Jonathan Korba, and Kumar Das. 2000. The 1999 DARPA off-line intrusion detection evaluation. *Computer networks* 34, 4 (2000), 579–595.

[55] Michael Xieyang Liu, Frederick Liu, Alexander J Fiannaca, Terry Koo, Lucas Dixon, Michael Terry, and Carrie J Cai. 2024. " We Need Structured Output": Towards User-centered Constraints on Large Language Model Output. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*. 1–9.

[56] Xinqian Liu, Jiadong Ren, Haitao He, Bing Zhang, Chen Song, and Yunxue Wang. 2021. A fast all-packets-based DDoS attack detection approach based on network graph and graph kernel. *Journal of Network and Computer Applications* 185 (2021), 103079. https://doi.org/10.1016/j.jnca.2021.103079

[57] Bernardo Louro, Raquel Abreu, Joana Cabral Costa, João B F. Sequeiros, and Pedro R M. Inácio. 2024. Analysis of the Capability and Training of Chat Bots in the Generation of Rules for Firewall or Intrusion Detection Systems. In *Proceedings of the 19th International Conference on Availability, Reliability and Security*. 1–7.

[58] Minzhao Lyu, Hassan Habibi Gharakheili, Craig Russell, and Vijay Sivaraman. 2021. Hierarchical anomaly-based detection of distributed DNS attacks on enterprise networks. *IEEE Transactions on Network and Service Management* 18, 1 (2021), 1031–1048.

[59] Jared Mathews, Prosenjit Chatterjee, Shankar Banik, and Cory Nance. 2022. A Deep Learning Approach to Create DNS Amplification Attacks. In *Proceedings of the 4th International Conference on Management Science and Industrial Engineering* (Chiang Mai, Thailand) *(MSIE '22)*. Association for Computing Machinery, New York, NY, USA, 429–435. https://doi.org/10.1145/3535782.3535838

[60] Kyle McGlynn, Hrishikesh B Acharya, and Minseok Kwon. 2019. Detecting BGP route anomalies with deep learning. In *IEEE INFOCOM 2019-IEEE Conference on Computer Communications Workshops (INFOCOM WKSHPS)*. IEEE, 1039–1040.

[61] Irom Lalit Meitei, Khundrakpam Johnson Singh, and Tanmay De. 2016. Detection of DDoS DNS Amplification Attack Using Classification Algorithm. In *Proceedings of the International Conference on Informatics and Analytics* (Pondicherry, India) *(ICIA-16)*. Association for Computing Machinery, New York, NY, USA, Article 81, 6 pages. https://doi.org/10.1145/2980258.2980431

[62] Jelena Mirkovic and Peter Reiher. 2004. A taxonomy of DDoS attack and DDoS defense mechanisms. *ACM SIGCOMM Computer Communication Review* 34, 2 (2004), 39–53.

[63] Yisroel Mirsky, Tomer Doitshman, Yuval Elovici, and Asaf Shabtai. 2018. Kitsune: An Ensemble of Autoencoders for Online Network Intrusion Detection. In *Network and Distributed Systems Security (NDSS) Symposium*.

[64] Nour Moustafa and Jill Slay. 2015. UNSW-NB15: a comprehensive data set for network intrusion detection systems (UNSW-NB15 network data set). In *2015 Military Communications and Information Systems Conference (MilCIS)*. IEEE, 1–6.

[65] Khan Muhammad, Mohammad S. Obaidat, Tanveer Hussain, Javier Del Ser, Neeraj Kumar, Mohammad Tanveer, and Faiyaz Doctor. 2021. Fuzzy Logic in Surveillance Big Video Data Analysis: Comprehensive Review, Challenges, and Research Directions. *ACM Comput. Surv.* 54, 3, Article 68 (May 2021), 33 pages. https://doi.org/10.1145/3444693

[66] Nura Shifa Musa, Nada Masood Mirza, Saida Hafsa Rafique, Amira Mahamat Abdallah, and Thangavel Murugan. 2024. Machine Learning and Deep Learning Techniques for Distributed Denial of Service Anomaly Detection in Software Defined Networks—Current Research Solutions. *IEEE Access* 12 (2024), 17982–18011. https://doi.org/10.1109/ACCESS.2024.3360868

[67] Ali Mustapha, Rida Khatoun, Sherali Zeadally, Fadlallah Chbib, Ahmad Fadlallah, Walid Fahs, and Ali El Attar. 2023. Detecting DDoS attacks using adversarial neural network. *Computers & Security* 127 (2023), 103117.

[68] Maryam M. Najafabadi, Taghi M. Khoshgoftaar, Chad Calvert, and Clifford Kemp. 2017. A Text Mining Approach for Anomaly Detection in Application Layer DDoS Attacks. *Proceedings of the Thirtieth International Florida Artificial Intelligence Research Society Conference* (2017).

[69] Maryam M. Najafabadi, Taghi M. Khoshgoftaar, Amri Napolitano, and Charles Wheelus. 2016. RUDY Attack: Detection at the Network Level and Its Important Features. *Proceedings of the Twenty-Ninth International Florida Artificial Intelligence Research Society Conference* (2016).

[70] Mohammad Najafimehr, Sajjad Zarifzadeh, and Seyedakbar Mostafavi. 2023. DDoS attacks and machine-learning-based detection methods: A survey and taxonomy. *Engineering Reports* 5, 12 (2023), e12697. https://doi.org/10.1002/eng2.12697 arXiv:https://onlinelibrary.wiley.com/doi/pdf/10.1002/eng2.12697

[71] Mohamed Nassar, Radu State, and Olivier Festor. 2008. Monitoring SIP Traffic Using Support Vector Machines. In *11th International Symposium on Recent Advances in Intrusion Detection*, Vol. 5230. 311–330. https://doi.org/10.1007/978-3-540-87403-4_17

[72] Janmenjoy Nayak, Bighnaraj Naik, and HSr Behera. 2015. Fuzzy C-means (FCM) clustering algorithm: a decade review from 2000 to 2014. In *Computational Intelligence in Data Mining-Volume 2: Proceedings of the International Conference on CIDM, 20-21 December 2014*. Springer, 133–149.

[73] Beny Nugraha, Naina Kulkarni, and Akash Gopikrishnan. 2021. Detecting adversarial DDoS attacks in software-defined networking using deep learning techniques and adversarial training. In *2021 IEEE International Conference on Cyber Security and Resilience (CSR)*. IEEE, 448–454.

[74] Ciprian Paduraru, Catalina Camelia Patilea, and Alin Stefanescu. 2024. CyberGuardian: An Interactive Assistant for Cybersecurity Specialists Using Large Language Models. *Proc. of ICSOFT* 24 (2024), 442–449.

[75] Amir Pakmehr, Andreas Aßmuth, Negar Taheri, and Ali Ghaffari. 2024. DDoS attack detection techniques in IoT networks: a survey. *Cluster Computing* 27, 10 (01 Dec 2024), 14637–14668. https://doi.org/10.1007/s10586-024-04662-6

[76] Sungho Park, Youngjun Kim, Hyungoo Choi, Yeunwoong Kyung, and Jinwoo Park. 2021. HTTP DDoS Flooding Attack Mitigation in Software-Defined Networking. *IEICE Transactions on Information and Systems* E104.D, 9 (2021), 1496–1499.

[77] Xiao Peng, Weiqing Huang, and Zhixin Shi. 2019. Adversarial attack against DoS intrusion detection: An improved boundary-based method. In *2019 IEEE 31st International Conference on Tools with Artificial Intelligence (ICTAI)*. IEEE, 1288–1295.

[78] Hung Viet Pham, Shangshu Qian, Jiannan Wang, Thibaud Lutellier, Jonathan Rosenthal, Lin Tan, Yaoliang Yu, and Nachiappan Nagappan. 2020. Problems and opportunities in training deep learning software systems: An analysis of variance. In *Proceedings of the 35th IEEE/ACM international conference on automated software engineering*. 771–783.

[79] Steven T Piantadosi. 2014. Zipf's word frequency law in natural language: A critical review and future directions. *Psychonomic bulletin & review* 21 (2014), 1112–1130.

[80] Md Abdul Quadir, J. Christy Jackson, J. Prassanna, K. Sathyarajasekaran, K. Kumar, H. Sabireen, Shivam Ubarhande, V. Vijaya Kumar, Vijayakumar Varadarajan, Piet Kommers, Vincenzo Piuri, and V. Subramaniyaswamy. 2020. An efficient algorithm to detect DDoS amplification attacks. *J. Intell. Fuzzy Syst.* 39, 6 (Jan. 2020), 8565–8572. https://doi.org/10.3233/JIFS-189173

[81] Riyadh Rahef Nuiaa, Selvakumar Manickam, and Ali Hakem ALsaeedi. 2022. A Comprehensive Review of DNS-based Distributed Reflection Denial of Service (DRDoS) Attacks: State-of-the-Art. *International Journal on Advanced Science, Engineering and Information Technology* 12, 6 (Dec. 2022), 2452–2461. https://doi.org/10.18517/ijaseit.12.6.17280

[82] Eguturi Manjith Kumar Reddy, Akash Gurrala, Vasireddy Bindu Hasitha, and Korupalli V Rajesh Kumar. 2022. Introduction to Naive Bayes and a review on its subtypes with applications. *Bayesian reasoning and gaussian processes for machine learning applications* (2022), 1–14.

[83] Vinicius Rios, Pedro Inacio, Damien Magoni, and Mario Freire. 2024. Detection of Slowloris Attacks using Machine Learning Algorithms. In *39th ACM/SIGAPP Symposium on Applied Computing (SAC 2024)*. ACM, 1321–1330.

[84] Christian Rossow. 2014. Amplification Hell: Revisiting Network Protocols for DDoS Abuse. In *Network and Distributed System Security Symposium*. 1–15.

[85] Mher Safaryan and Peter Richtárik. 2021. Stochastic sign descent methods: New algorithms and better theory. In *International Conference on Machine Learning*. PMLR, 9224–9234.

[86] Mahmoud Said Elsayed, Nhien-An Le-Khac, Soumyabrata Dev, and Anca Delia Jurcut. 2020. Network anomaly detection using LSTM based autoencoder. In *Proceedings of the 16th ACM Symposium on QoS and Security for Wireless and Mobile Networks*. 37–45.

[87] Samed Saka, Ali Al-Ataby, and Valerio Selis. 2023. Generating Synthetic Tabular Data for DDoS Detection Using Generative Models. In *Proceedings of TrustCom*. 1436–1442. https://doi.org/10.1109/TrustCom60117.2023.00196

[88] Mohammad A Salahuddin, Vahid Pourahmadi, Hyame Assem Alameddine, Md Faizul Bari, and Raouf Boutaba. 2021. Chronos: DDoS attack detection using time-based autoencoder. *IEEE Transactions on Network and Service Management* 19, 1 (2021), 627–641.

[89] Vitalii Savchenko, Oleh Ilin, Nikolay Hnidenko, Olga Tkachenko, Oleksander Laptiev, and Svitlana Lehominova. 2020. Detection of Slow DDoS Attacks based on User's Behavior Forecasting. *International Journal of Emerging Trends in Engineering Research* 8, 5 (2020).

[90] Naeem Seliya, Azadeh Abdollah Zadeh, and Taghi M Khoshgoftaar. 2021. A literature review on one-class classification and its potential applications in big data. *Journal of Big Data* 8 (2021), 1–31.

[91] Malliga Senthil, P. Nandhini, and Sv Kogilavani. 2022. A Comprehensive Review of Deep Learning Techniques for the Detection of (Distributed) Denial of Service Attacks. *Information Technology and Control* 51 (03 2022), 180–215. https://doi.org/10.5755/j01.itc.51.1.29595

[92] Iman Sharafaldin, Arash Habibi Lashkari, Ali A Ghorbani, et al. 2018. Toward generating a new intrusion detection dataset and intrusion traffic characterization. *ICISSP* 1 (2018), 108–116.

[93] Iman Sharafaldin, Arash Habibi Lashkari, Saqib Hakak, and Ali A Ghorbani. 2019. Developing realistic distributed denial of service (DDoS) attack dataset and taxonomy. In *2019 International Carnahan Conference on Security Technology (ICCST)*. IEEE, 1–8.

[94] Chin-Shiuh Shieh, Wan-Wei Lin, Thanh-Tuan Nguyen, Chi-Hong Chen, Mong-Fong Horng, and Denis Miu. 2021. Detection of unknown DDoS attacks with deep learning and Gaussian mixture model. *Applied Sciences* 11, 11 (2021), 5213.

[95] Jugal Shroff, Rahee Walambe, Sunil Kumar Singh, and Ketan Kotecha. 2022. Enhanced security against volumetric DDoS attacks using adversarial machine learning. *Wireless Communications and Mobile Computing* 2022, 1 (2022), 5757164.

[96] Mohammad Shurman, Rami Khrais, Abdulrahman Yateem, et al. 2020. DoS and DDoS attack detection using deep learning and IDS. *Int. Arab J. Inf. Technol* 17, 4A (2020), 655–661.

[97] Anshuman Singh and Brij B. Gupta. 2022. Distributed Denial-of-Service (DDoS) Attacks and Defense Mechanisms in Various Web-Enabled Computing Platforms: Issues, Challenges, and Future Research Directions. *International Journal on Semantic Web and Information Systems (IJSWIS)* 18, 1 (Jan. 2022), 1–43. https://doi.org/10.4018/IJSWIS.297143

[98] Yinghao Su, Dapeng Xiong, Kechang Qian, and Yu Wang. 2024. A Comprehensive Survey of Distributed Denial of Service Detection and Mitigation Technologies in Software-Defined Network. *Electronics* 13, 4 (2024). https://doi.org/10.3390/electronics13040807

[99] Saowanit Sukparungsee, Yupaporn Areepong, and Rattikarn Taboran. 2020. Exponentially weighted moving average—Moving average charts for monitoring the process mean. *Plos one* 15, 2 (2020), e0228208.

[100] S. Mary Saira Bhanu T. Raja Sree. 2019. Detection of HTTP flooding attacks in cloud using fuzzy bat clustering. *Neural Computing and Applications* 32 (2019), 9603–9619.

[101] Yang Xiao Tasnuva Mahjabin. 2019. DNS Flood Attack Mitigation Utilizing Hot-Lists and Stale Content Updates. *Security, Privacy, and Anonymity in Computation, Communication, and Storage* (2019), 289–296.

[102] Mahbod Tavallaee, Ebrahim Bagheri, Wei Lu, and Ali A Ghorbani. 2009. A detailed analysis of the KDD CUP 99 data set. In *2009 IEEE symposium on computational intelligence for security and defense applications.* Ieee, 1–6.

[103] Nikhil Tripathi and Neminath Hubballi. 2018. Detecting Stealth DHCP Starvation Attack using Machine Learning Approach. *Journal of Computer Virology and Hacking Techniques* 14 (08 2018). https://doi.org/10.1007/s11416-017-0310-x

[104] Nikhil Tripathi and Neminath Hubballi. 2021. Application layer denial-of-service attacks and defense mechanisms: a survey. *ACM Computing Surveys (CSUR)* 54, 4 (2021), 1–33.

[105] Eric Tzeng, Judy Hoffman, Kate Saenko, and Trevor Darrell. 2017. Adversarial Discriminative Domain Adaptation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition.* 7167–7176.

[106] Greg Van Houdt, Carlos Mosquera, and Gonzalo Nápoles. 2020. A review on the long short-term memory model. *Artificial Intelligence Review* 53, 8 (2020), 5929–5955.

[107] Saurabh Verma, Ali Hamieh, Jun Ho Huh, Henrik Holm, Siva Raj Rajagopalan, Maciej Korczynski, and Nina Fefferman. 2016. Stopping amplified DNS DDoS attacks through distributed query rate sharing. In *2016 11th International Conference on Availability, Reliability and Security (ARES).* IEEE, 69–78.

[108] Abdullahi Aishatu Wabi, Ismaila Idris, Olayemi Mikail Olaniyi, and Joseph A. Ojeniyi. 2024. DDOS attack detection in SDN: Method of attacks, detection techniques, challenges and research gaps. *Computers & Security* 139 (2024), 103652. https://doi.org/10.1016/j.cose.2023.103652

[109] Daniel Wagner, Daniel Kopp, Matthias Wichtlhuber, Christoph Dietzel, Oliver Hohlfeld, Georgios Smaragdakis, and Anja Feldmann. 2021. United we stand: Collaborative detection and mitigation of amplification DDoS attacks at scale. In *Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security.* 970–987.

[110] Jin Wang, XiaoLong Yang, KePing Long Min Zhang, and Jie Xu. 2014. HTTP-SoLDiER: An HTTP-flooding attack detection scheme with the large deviation principle. *Science China Information Sciences* 57 (2014), 1–15.

[111] Tongze Wang, Xiaohui Xie, Lei Zhang, Chuyi Wang, Liang Zhang, and Yong Cui. 2024. ShieldGPT: An LLM-based framework for DDoS mitigation. In *Proceedings of the 8th Asia-Pacific Workshop on Networking.* 108–114.

[112] Zheng Wang. 2012. Analysis of Flooding DoS Attacks Utilizing DNS Name Error Queries. *KSII TRANSACTIONS ON INTERNET AND INFORMATION SYSTEMS* 6, 10 (2012).

[113] Yuanyuan Wei, Julian Jang-Jaccard, Fariza Sabrina, Amardeep Singh, Wen Xu, and Seyit Camtepe. 2021. AE-MLP: A hybrid deep learning approach for DDoS detection and classification. *IEEE Access* 9 (2021), 146810–146821.

[114] Yuanyuan Wei, Julian Jang-Jaccard, Fariza Sabrina, Wen Xu, Seyit Camtepe, and Aeryn Dunmore. 2023. Reconstruction-based LSTM-autoencoder for anomaly-based DDoS attack detection over multivariate time-series data. *arXiv preprint arXiv:2305.09475* (2023).

[115] Rey Wiyatno and Anqi Xu. 2018. Maximal jacobian-based saliency map attack. *arXiv preprint arXiv:1808.07945* (2018).

[116] Andrew Wei-Loong Wong, Say Leng Goh, Mohammad Kamrul Hasan, and Salmah Fattah. 2024. Multi-Hop and Mesh for LoRa Networks: Recent Advancements, Issues, and Recommended Applications. *ACM Comput. Surv.* 56, 6, Article 136 (Jan. 2024), 43 pages. https://doi.org/10.1145/3638241

[117] Hongzuo Xu, Guansong Pang, Yijie Wang, and Yongjun Wang. 2023. Deep isolation forest for anomaly detection. *IEEE Transactions on Knowledge and Data Engineering* 35, 12 (2023), 12591–12604.

[118] Lei Xu, Maria Skoularidou, Alfredo Cuesta-Infante, and Kalyan Veeramachaneni. 2019. Modeling tabular data using conditional gan. *Advances in neural information processing systems* 32 (2019).

[119] Zhenyu Yin, Shang Liu, and Guangyuan Xu. 2024. DrLLM: Prompt-Enhanced Distributed Denial-of-Service Resistance Method with Large Language Models. *arXiv preprint arXiv:2409.10561* (2024).

[120] Zihan Yu, Liang He, Zhen Wu, Xinyu Dai, and Jiajun Chen. 2023. Towards better chain-of-thought prompting strategies: A survey. *arXiv preprint arXiv:2310.04959* (2023).

[121] Martin Zadnik and Elena Carasec. 2023. AI infers DoS mitigation rules. *Journal of Intelligent Information Systems* 60, 2 (2023), 305–324.

[122] Chaoyun Zhang, Xavier Costa-Pérez, and Paul Patras. 2020. Tiki-taka: Attacking and defending deep learning-based intrusion detection systems. In *Proceedings of the 2020 ACM SIGSAC Conference on Cloud Computing Security Workshop.* 27–39.

[123] Chenhan Zhang, Shui Yu, Zhiyi Tian, and James J. Q. Yu. 2023. Generative Adversarial Networks: A Survey on Attack and Defense Perspective. *ACM Comput. Surv.* 56, 4, Article 91 (Nov. 2023), 35 pages. https://doi.org/10.1145/3615336

[124] Zhirui Zhang, Shujie Liu, Mu Li, Ming Zhou, and Enhong Chen. 2018. Bidirectional generative adversarial networks for neural machine translation. In *Proceedings of the 22nd conference on computational natural language learning.* 190–199.

[125] Ziming Zhao, Zhaoxuan Li, Zhihao Zhou, Jiongchi Yu, Zhuoxue Song, Xiaofei Xie, Fan Zhang, and Rui Zhang. 2024. DDoS family: A novel perspective for massive types of DDoS attacks. *Computers & Security* 138 (March 2024), 103663. https://doi.org/10.1016/j.cose.2023.103663

[126] Junhao Zheng, Shengjie Qiu, Chengming Shi, and Qianli Ma. 2025. Towards Lifelong Learning of Large Language Models: A Survey. *ACM Comput. Surv.* 57, 8, Article 193 (March 2025), 35 pages. https://doi.org/10.1145/3716629

[127] Bo Zong, Qi Song, Martin Renqiang Min, Wei Cheng, Cristian Lumezanu, Daeki Cho, and Haifeng Chen. 2018. Deep autoencoding gaussian mixture model for unsupervised anomaly detection. In *International conference on learning representations.*